

Lecture: Pixels and Filters

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:
Forsyth and Ponce, Computer Vision, Chapter 7

What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

Types of Images

Binary



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Types of Images

Binary



Gray Scale



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Types of Images

Binary



Gray Scale

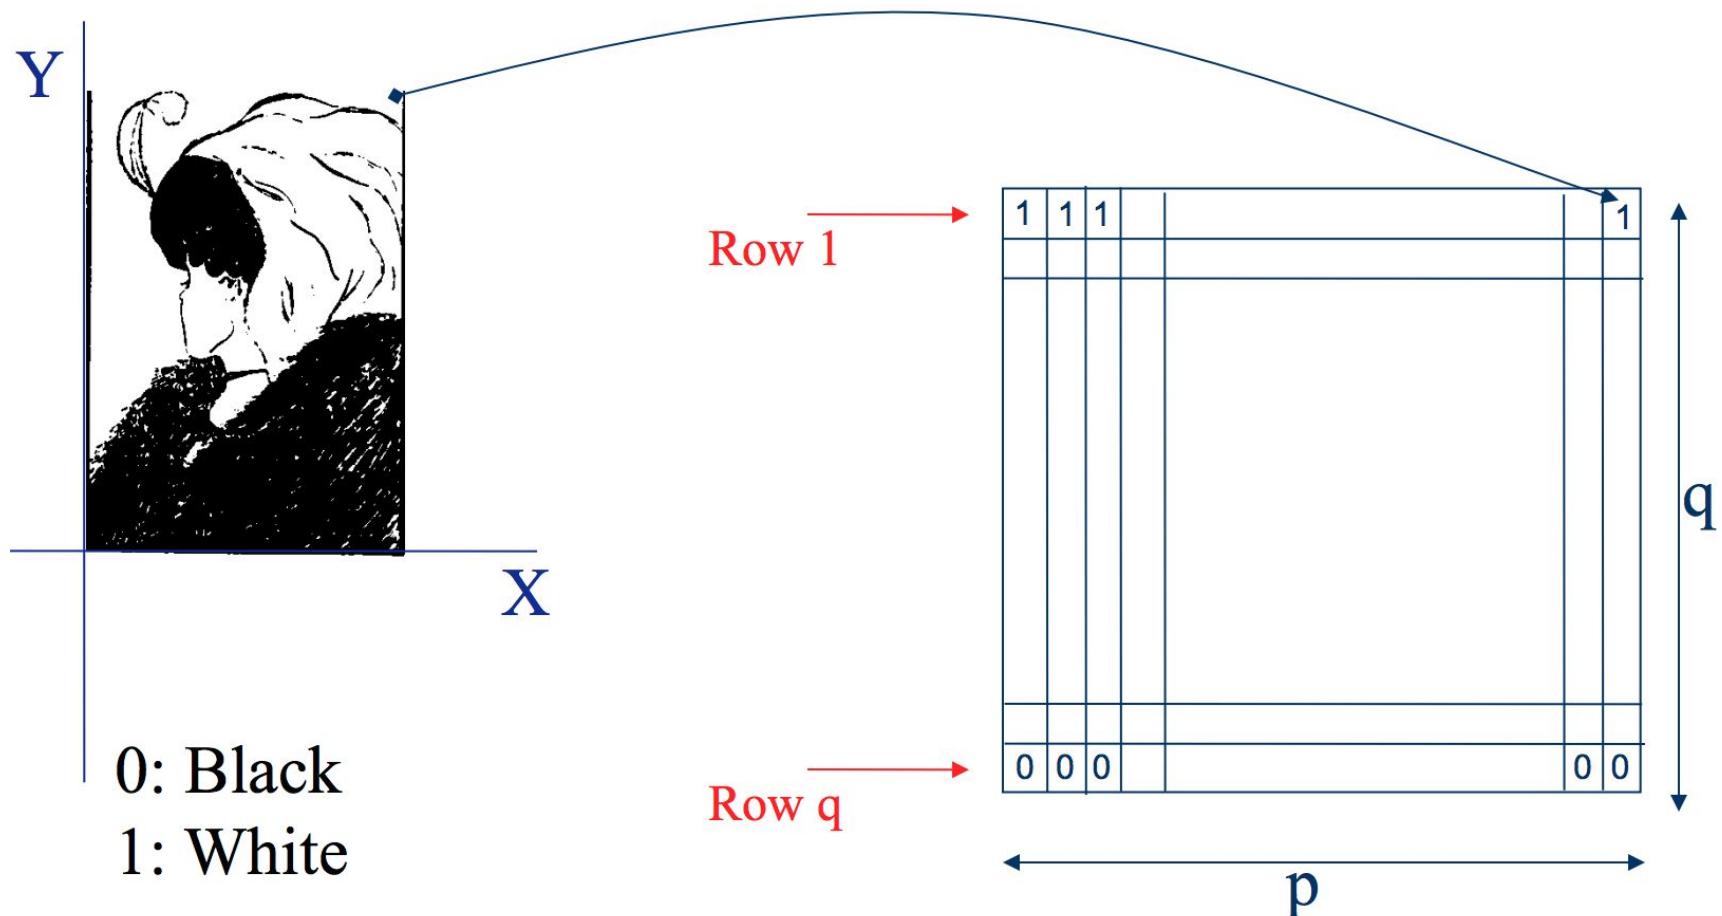


Color



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

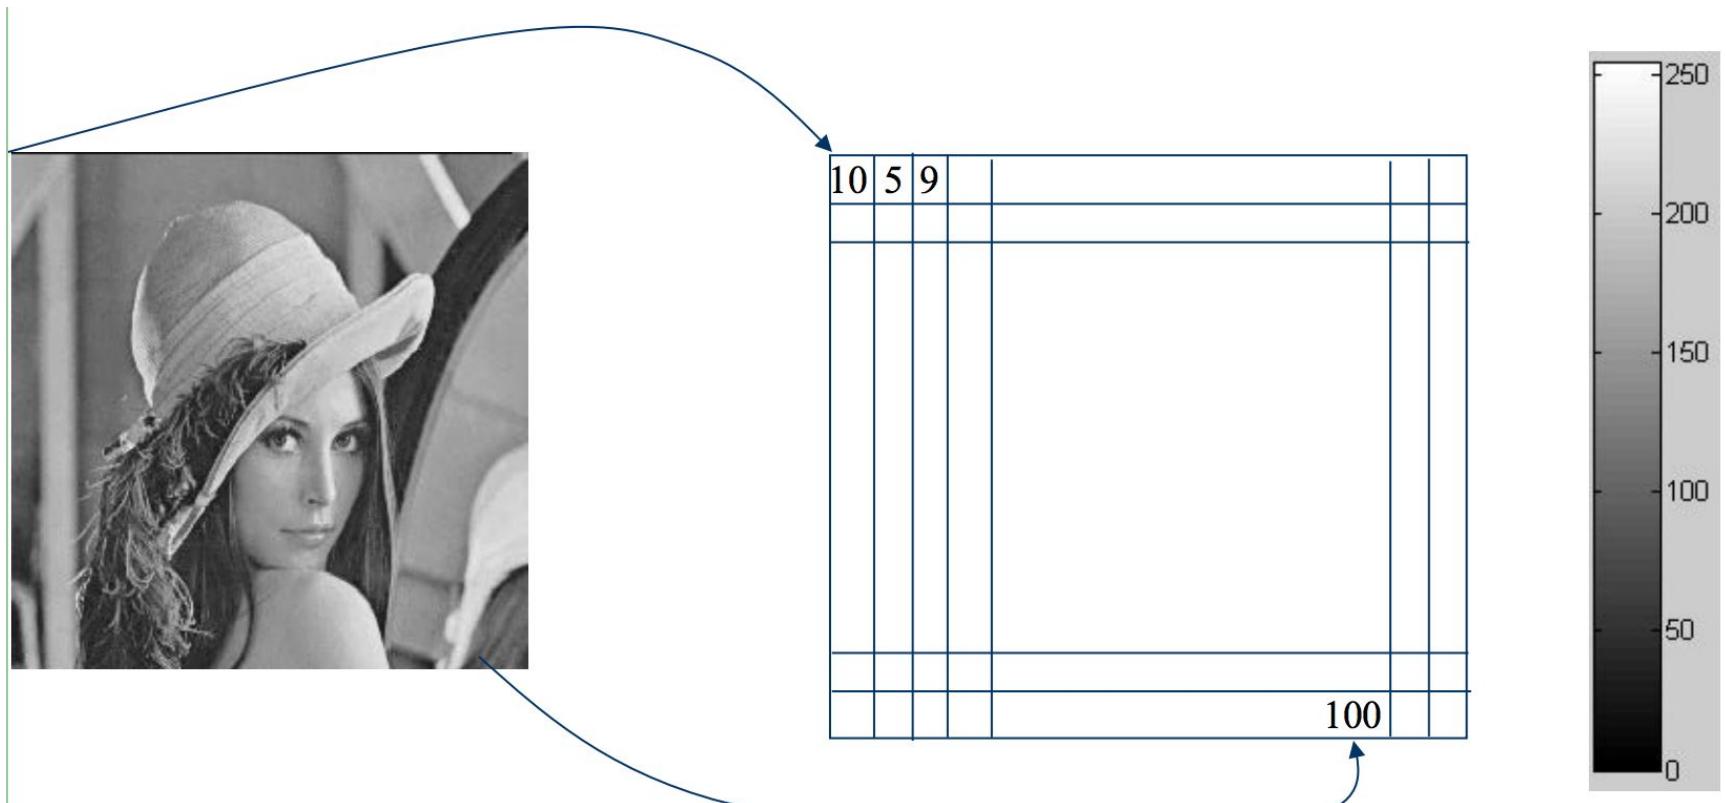
Binary image representation



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Ulas Bagci

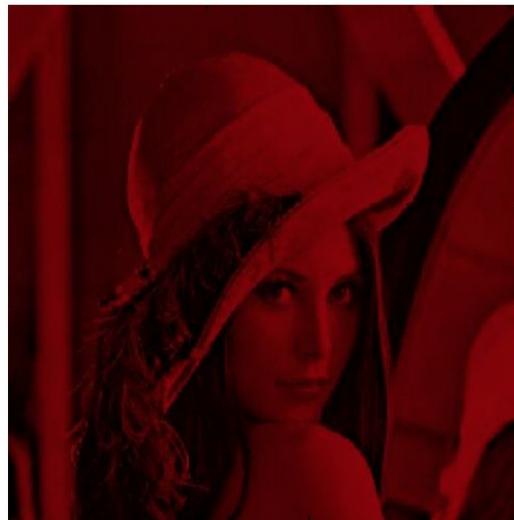
Grayscale image representation



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Ulas Bagci

Color Image - one channel



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Ulas Bagci

Color image representation

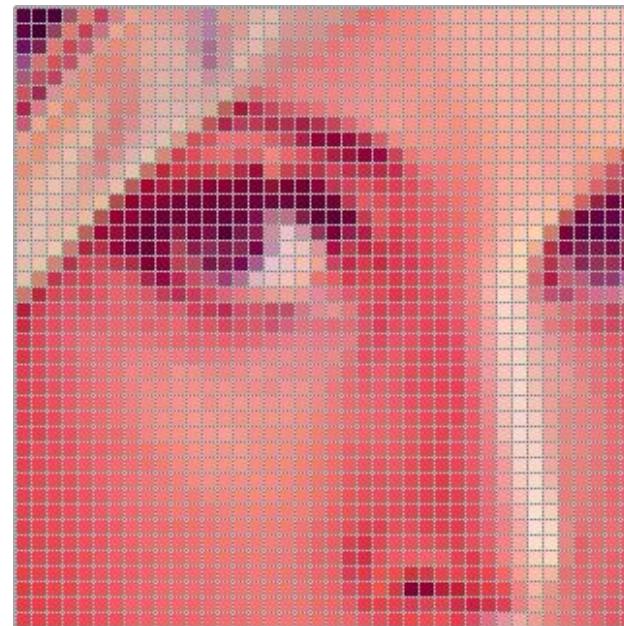
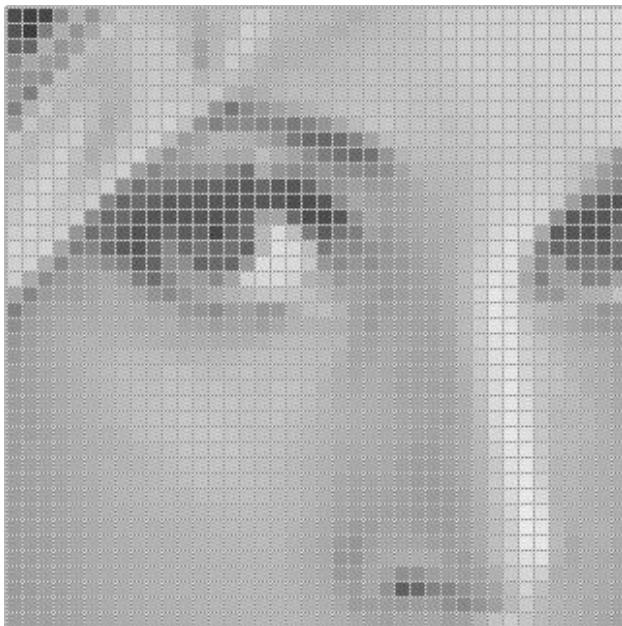


Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Ulas Bagci

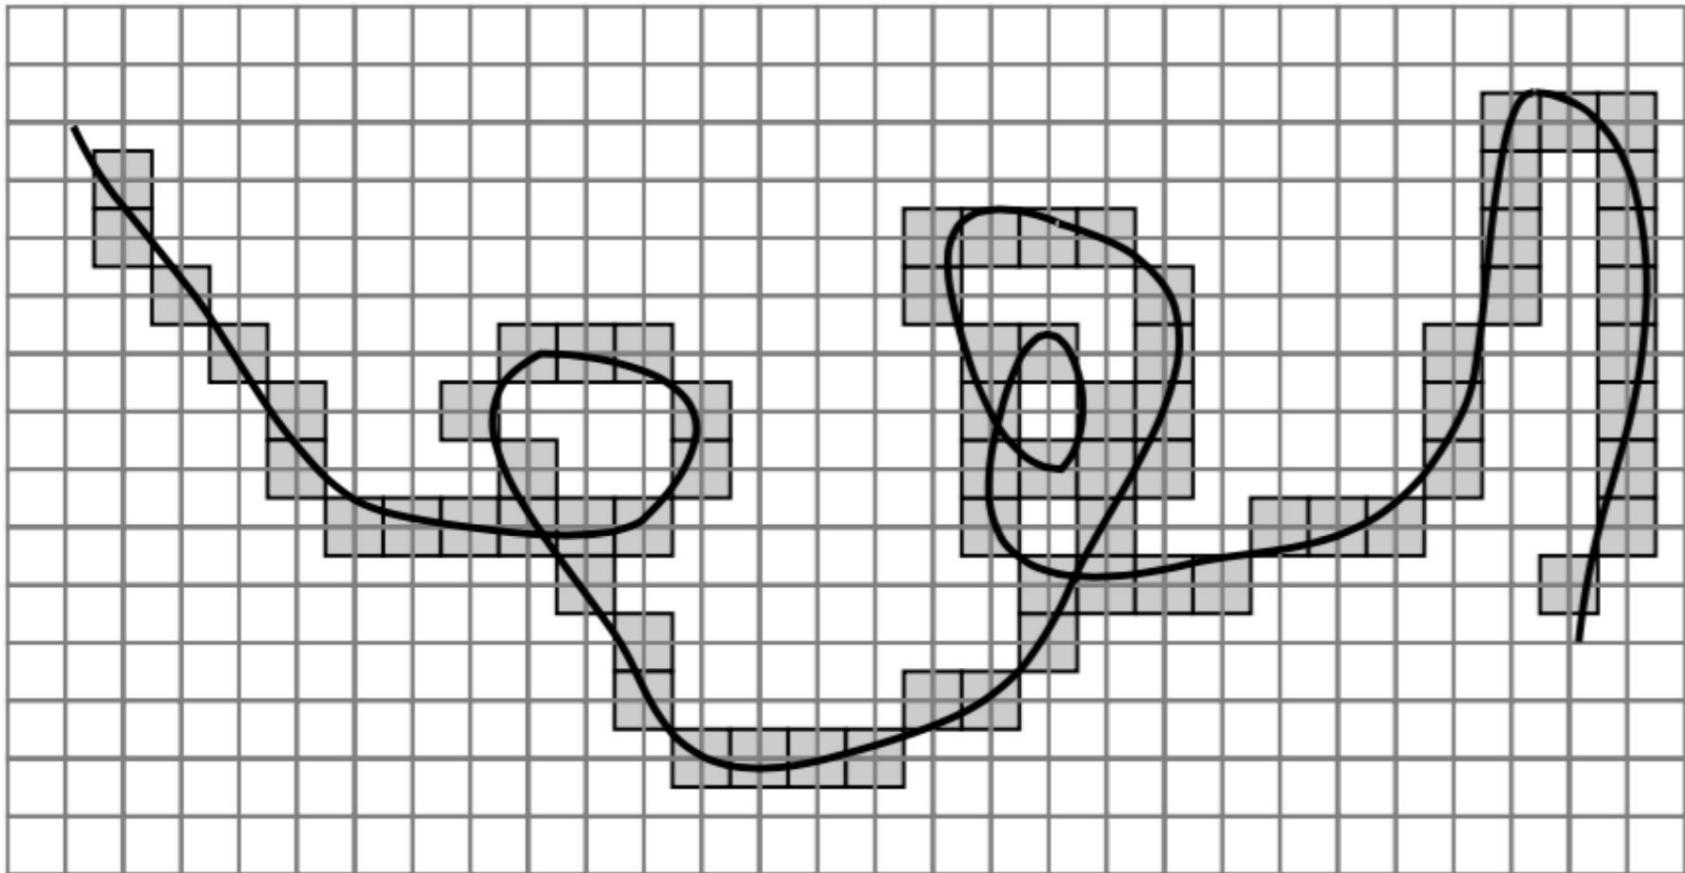
Images are sampled

What happens when we zoom into the images we capture?



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Errors due Sampling



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Ulas Bagci

Resolution

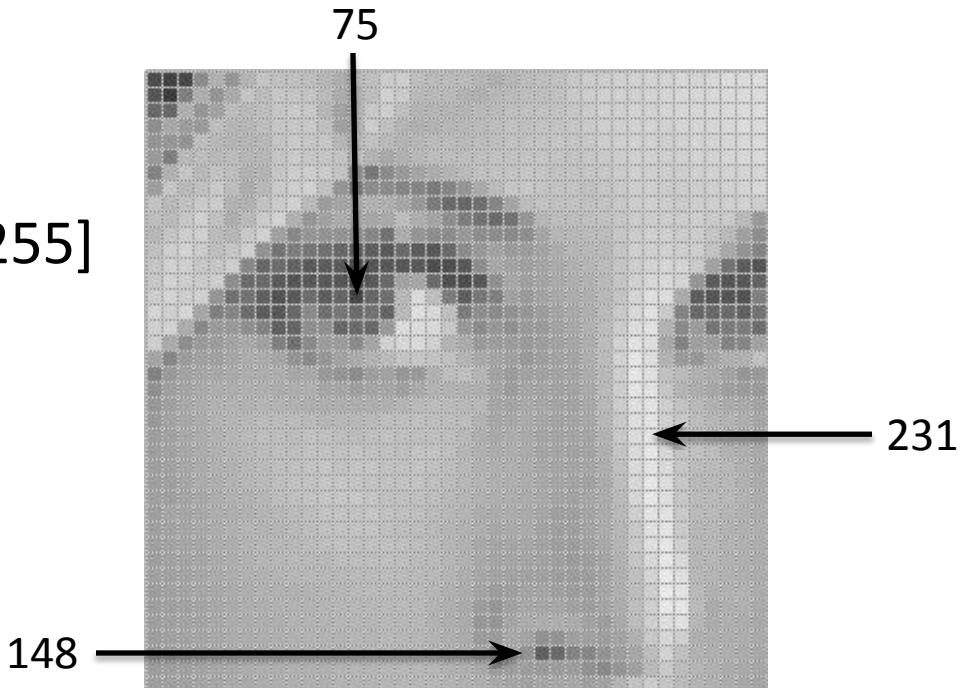
is a **sampling** parameter, defined in dots per inch (DPI) or equivalent measures of spatial pixel density

- Resolution of the camera and the resolution of the screen can both affect the quality.



Images are Sampled and Quantized

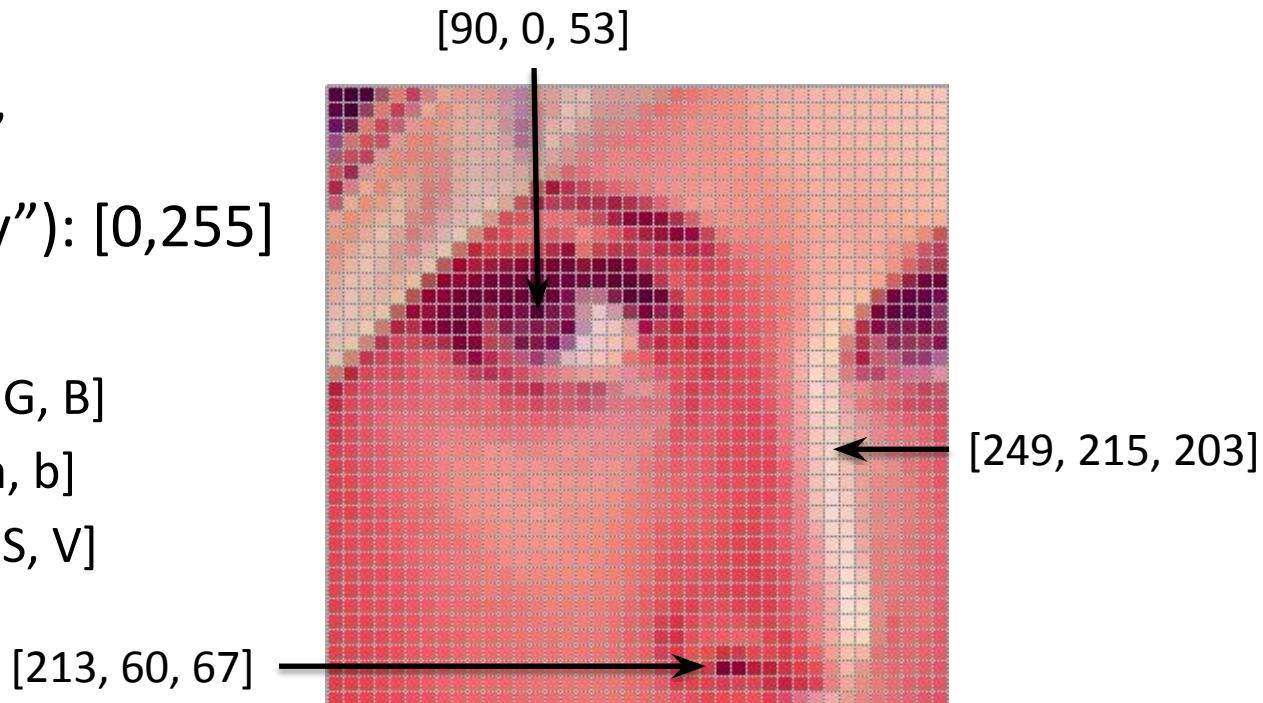
- An image contains discrete number of pixels
 - A simple example
 - Pixel value:
 - “grayscale”
- (or “intensity”): [0,255]



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Images are Sampled and Quantized

- An image contains discrete number of pixels
 - A simple example
 - Pixel value:
 - “grayscale”
(or “intensity”): [0,255]
 - “color”
 - RGB: [R, G, B]
 - Lab: [L, a, b]
 - HSV: [H, S, V]



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

With this loss of information (from sampling and quantization),

Can we still use images for useful tasks?

What we will learn today?

- Image sampling and quantization
- **Image histograms**
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:

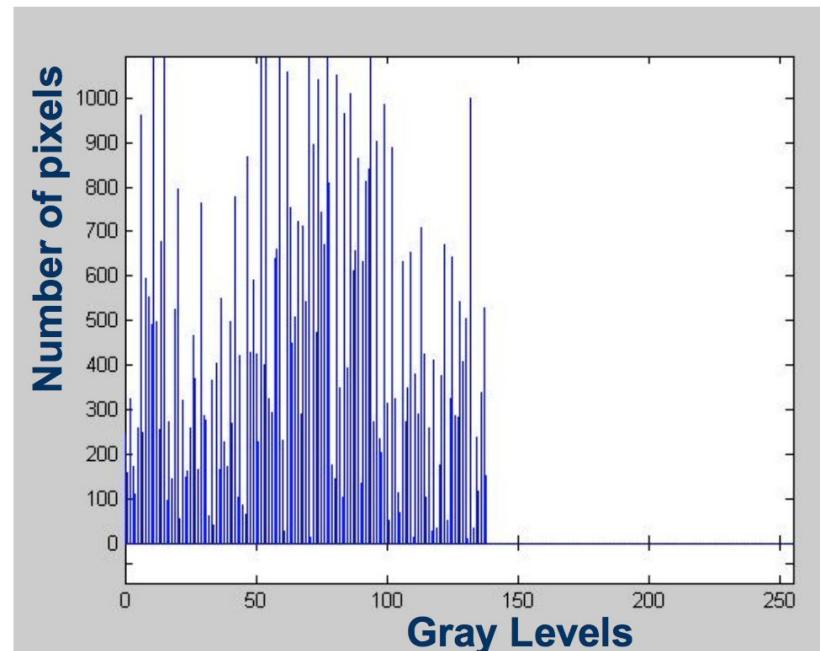
Forsyth and Ponce, Computer Vision, Chapter 7

Histogram

- Histogram captures the distribution of "countable elements" in the image, e.g. how frequently each gray level occurs in the image

Histogram

- Grayscale image - intensity histogram: one of the simplest *image representations*



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

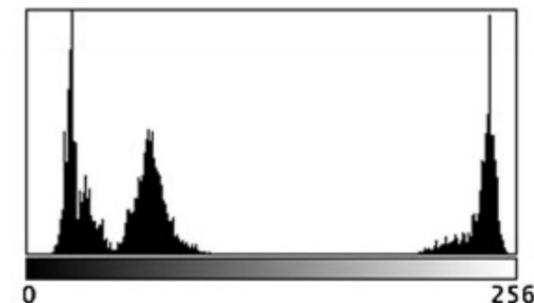
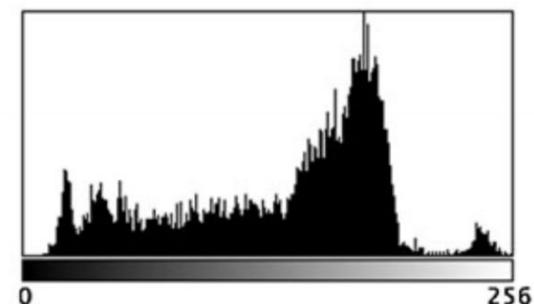
Histogram

Example implementation of the most basic grayscale histogram:

```
def histogram(im):
    h = np.zeros(256)
    for row in range(im.shape[0]):
        for col in range(im.shape[1]):
            val = im[row, col]
            h[val] += 1
```

Note: Python for loops can be slow, a faster implementation: see *numpy.histogram*.

Histogram



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Dr. Mubarak Shah

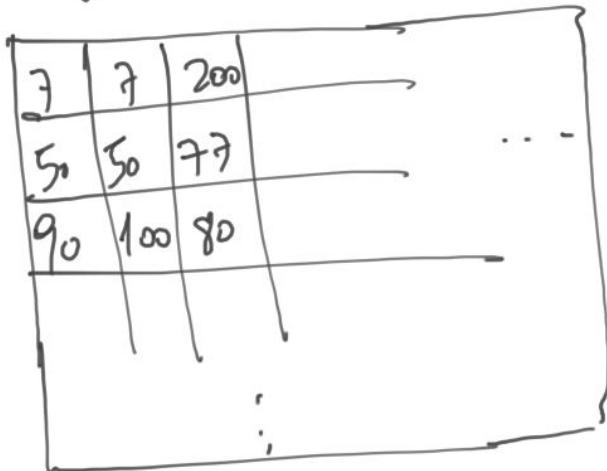
Histogram

7	7	200
50	50	77
90	100	80

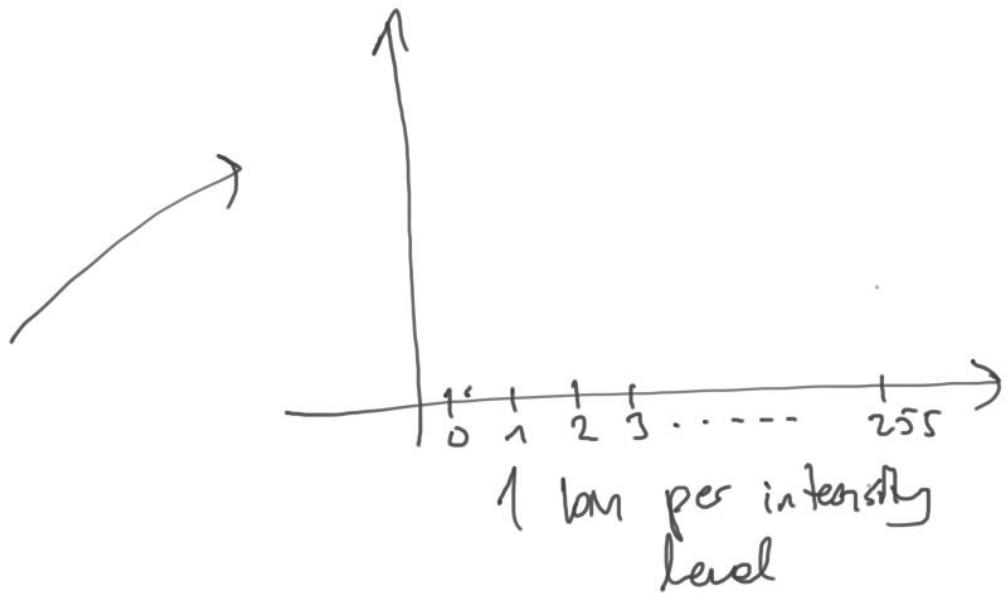
gray scale image

?

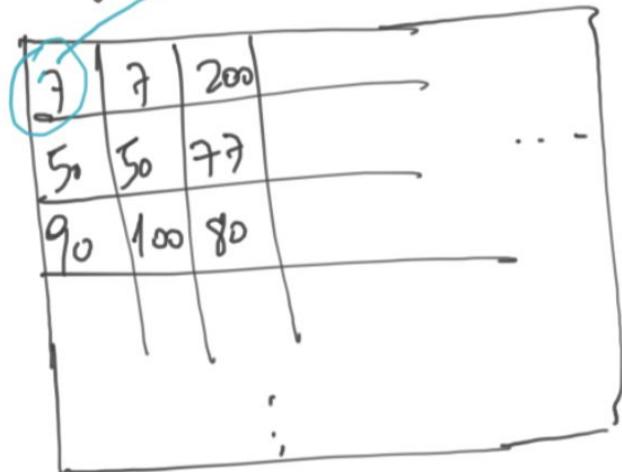
Histogram



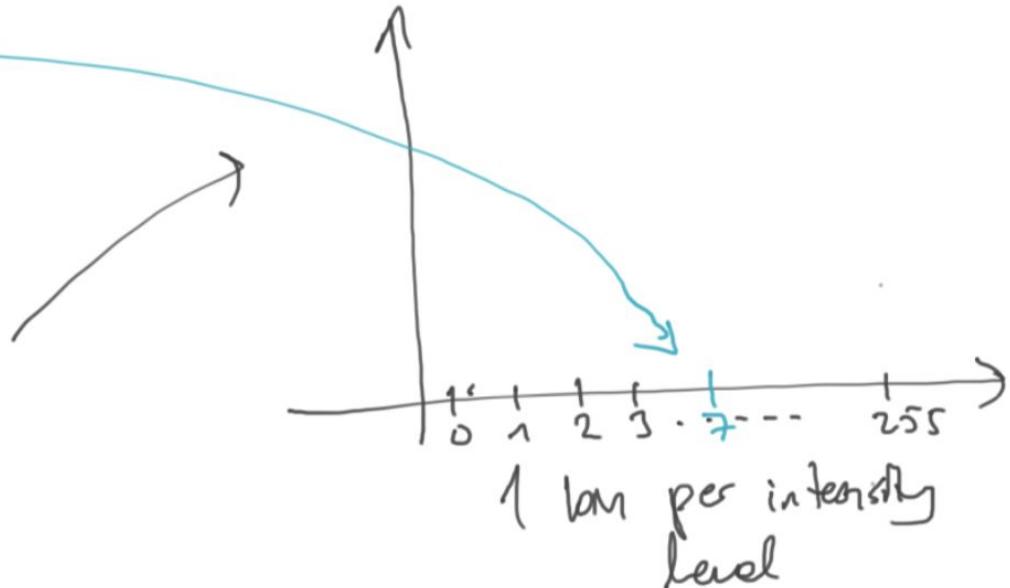
gray scale image



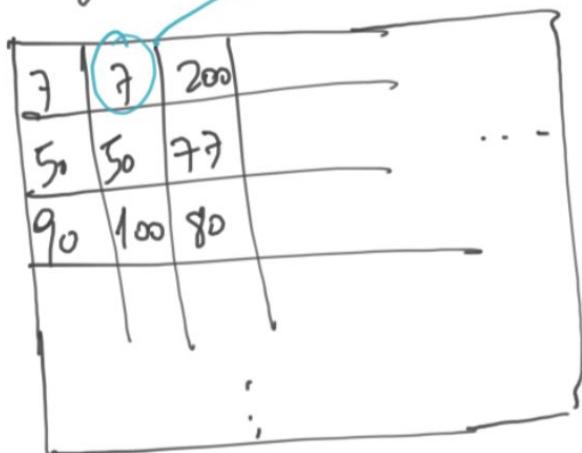
Histogram



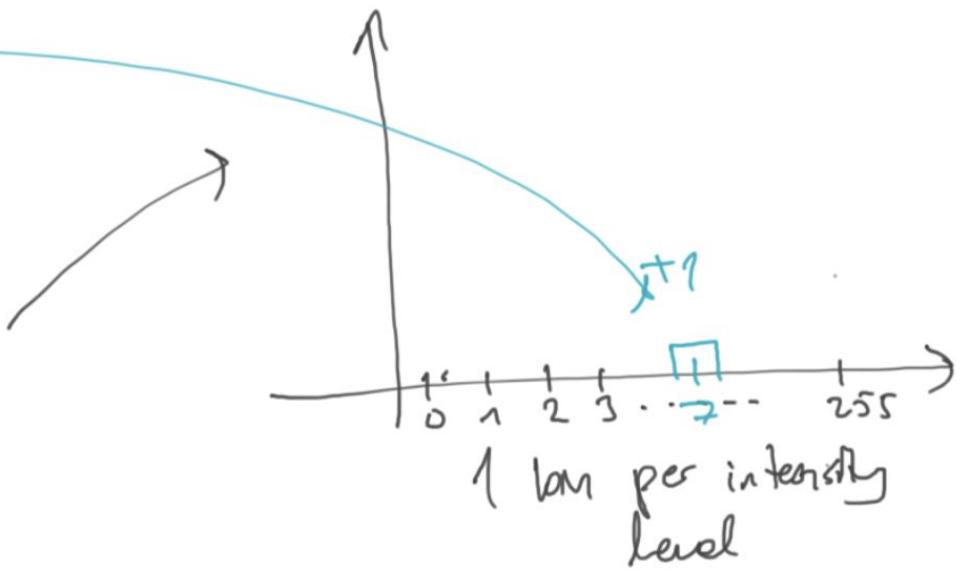
gray scale image



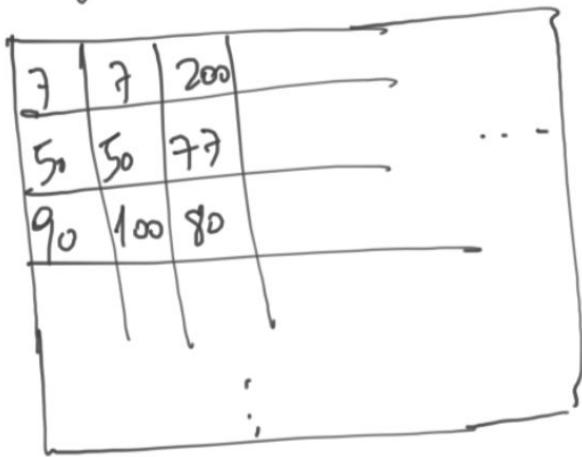
Histogram



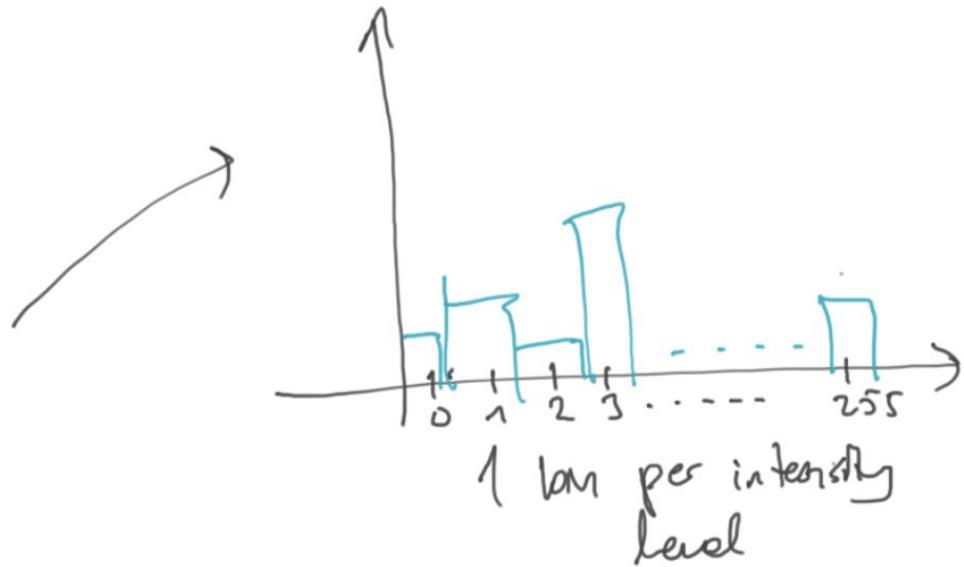
gray scale image



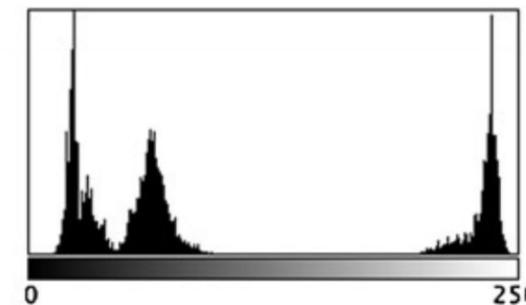
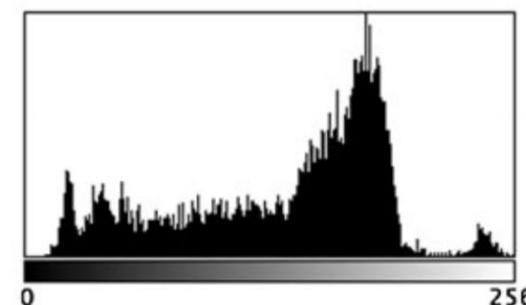
Histogram



gray scale image



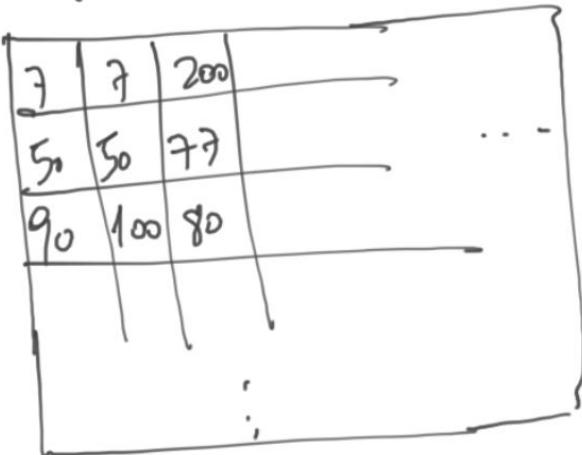
How to improve the histogram? (ie. What can go wrong with 1 bin per intensity level?)



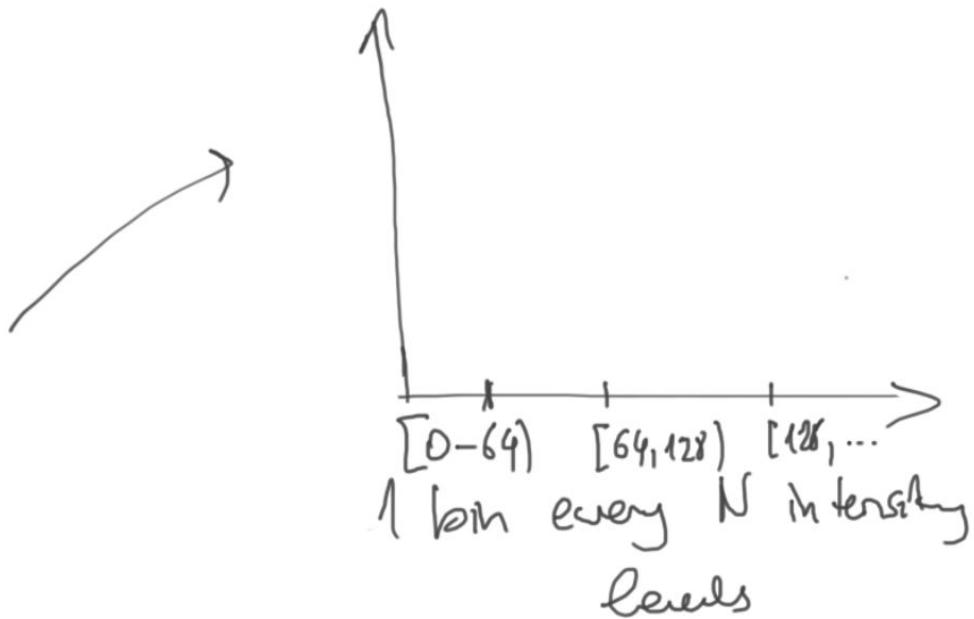
Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Dr. Mubarak Shah

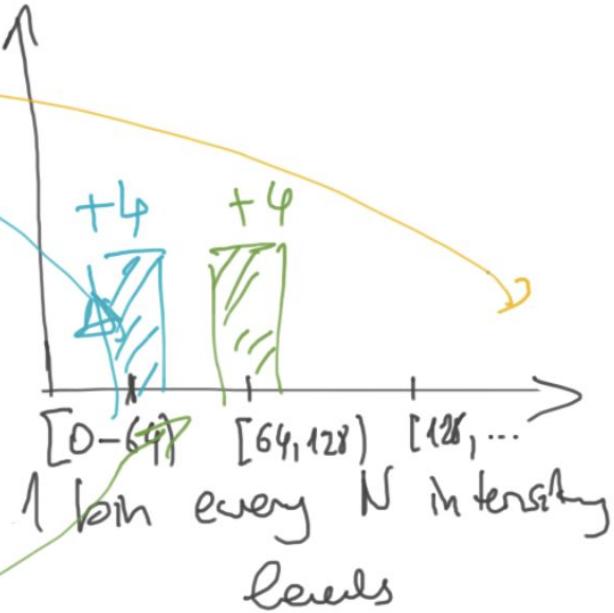
Histogram



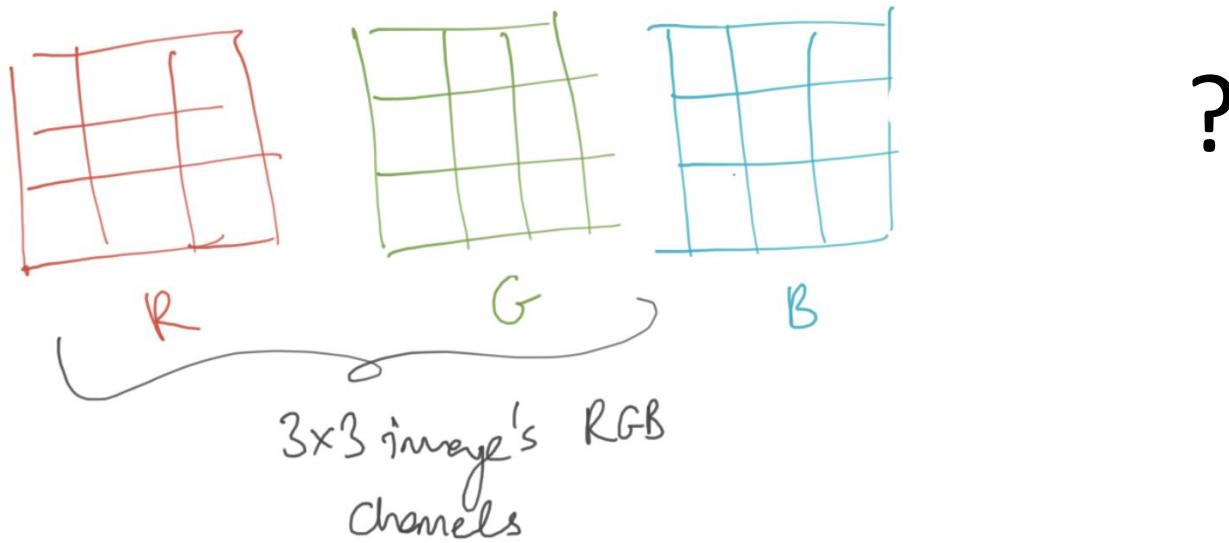
gray scale image



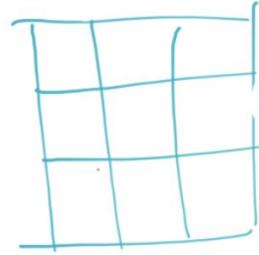
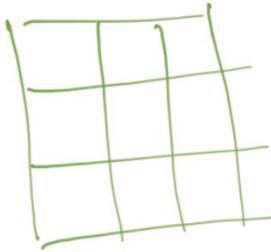
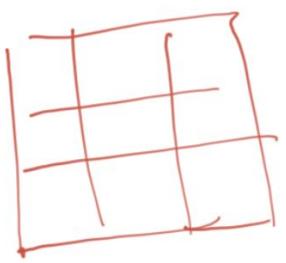
Histogram



Color histogram

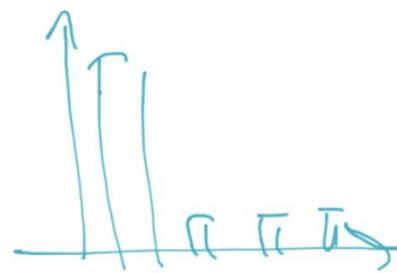
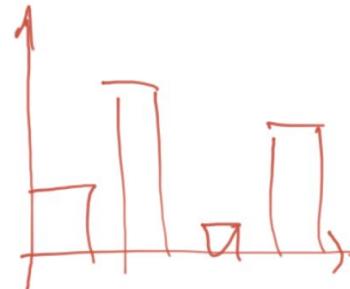


Color histogram

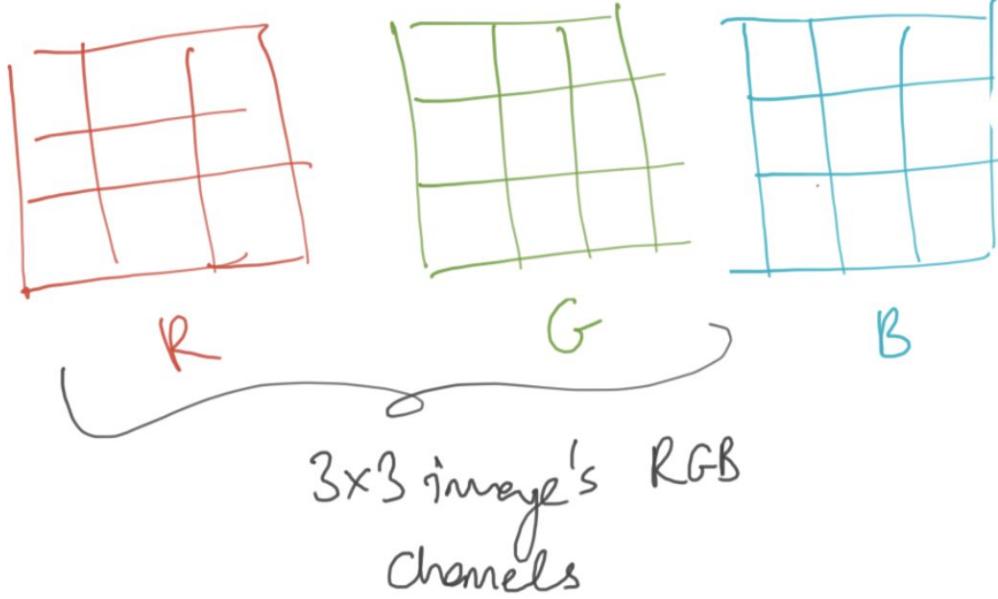


3x3 image's RGB
channels

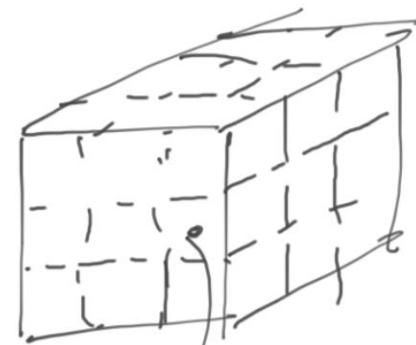
1 histogram per channel



Color histogram



Joint histogram



A single cell J :

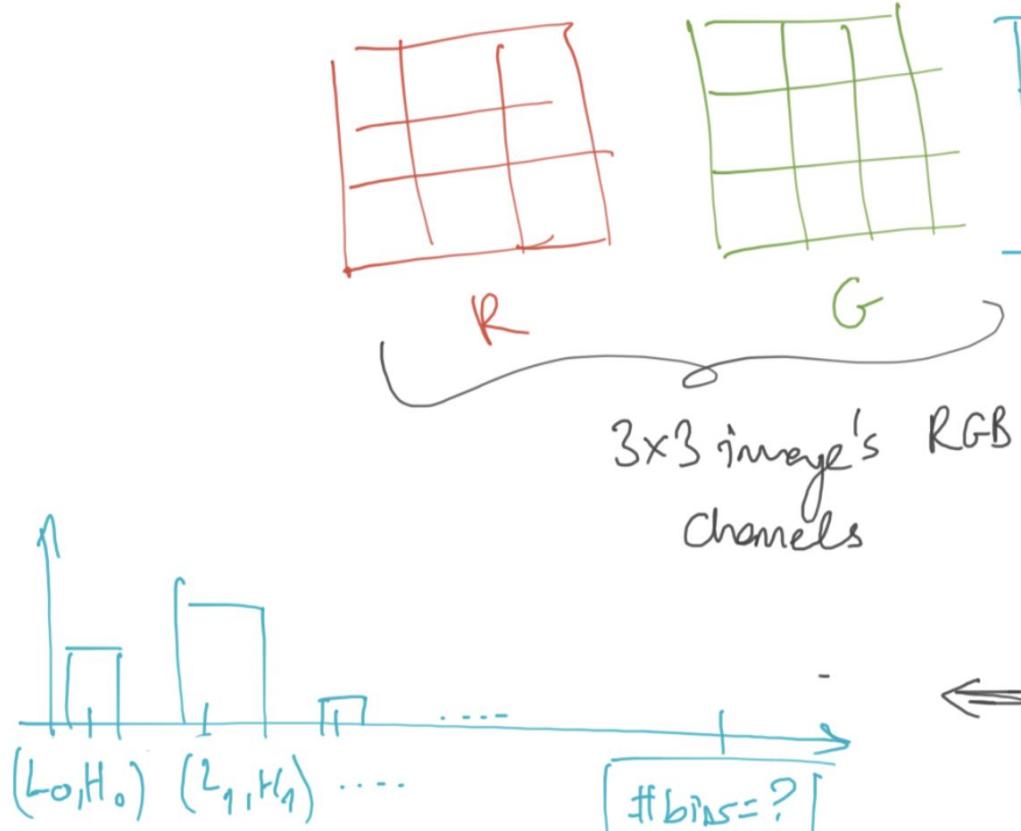
Count of pixels in the range:

$$L_J^r \leq r < H_J^r, L_J^g \leq g < H_J^g$$

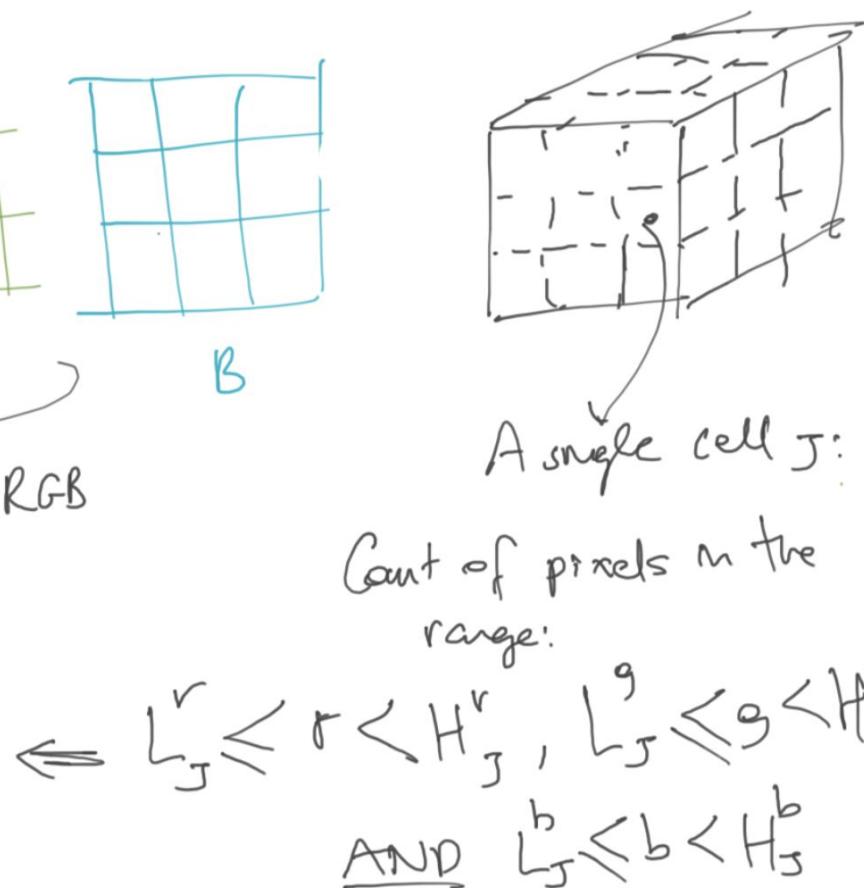
AND

$$L_J^b \leq b < H_J^b$$

Color histogram



Joint histogram



Color histogram

THINK OF ADVANTAGES/DISADVANTAGES:

- Number of bins in grayscale/color histograms
- Grayscale vs color histogram
- Per-channel vs joint (n -dim) color histogram



#bins=?

Joint histogram



3×3 image's
Channels

Count of pixels in the
range:

$$\Leftrightarrow L_J^r \leq r < H_J^r, L_J^g \leq g < H_J^g$$

$$\text{AND } L_J^b \leq b < H_J^b$$

n bins
per chart

$$d_{\frac{1}{2}} \quad d_{30}$$

$$\boxed{d = 60}$$

$$4 \quad 12 \quad 64$$

$$8 \quad 24 \quad 512$$

$$10 \quad 30 \quad 1000$$

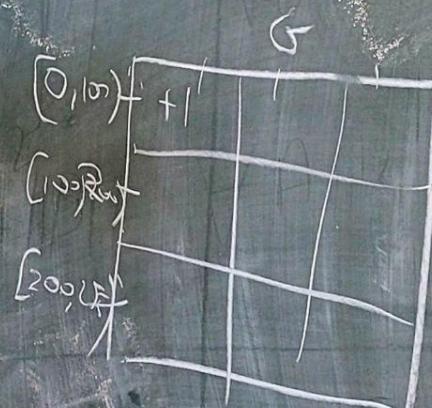
(20)

$$30 \quad 90 \quad 2700$$

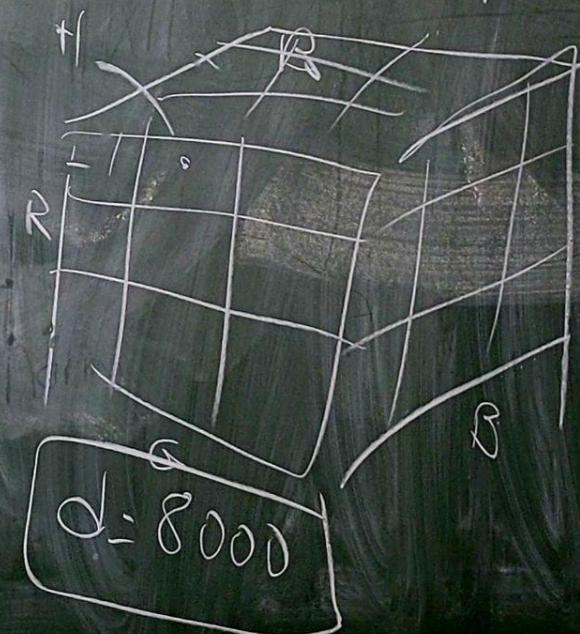
$$150 \quad 450 \quad -$$

R	G	B
0	0	255
0	255	0
0	0	255
0	255	0

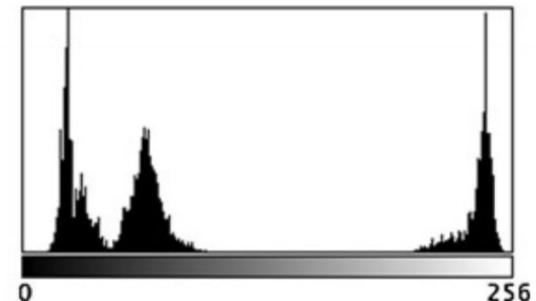
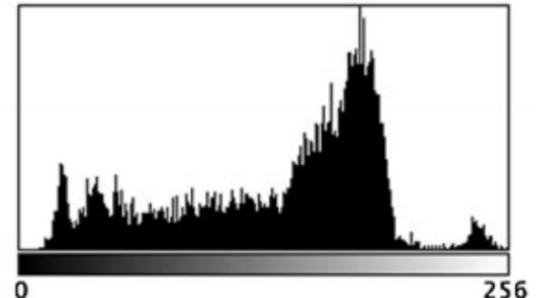
R	G	B
0	255	255
0	255	255
0	0	0
0	0	0



RC → consistent
RB / BG



Re-think of histogram's application



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

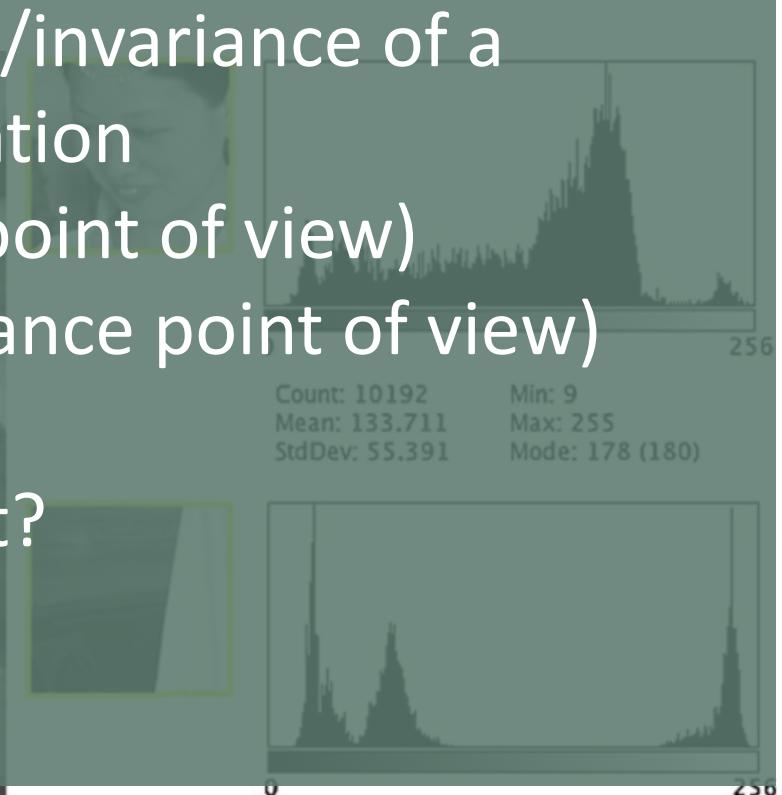
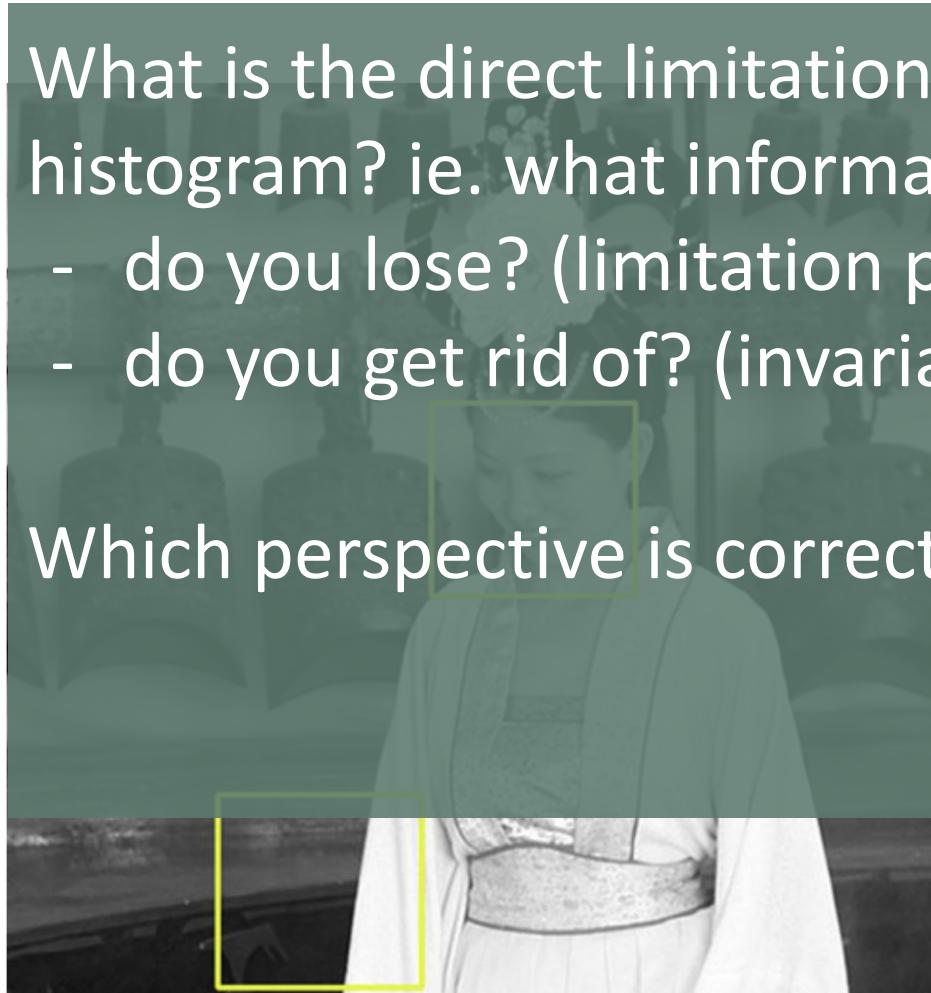
Slide credit: Dr. Mubarak Shah

Re-think of histogram's application

What is the direct limitation/invariance of a histogram? ie. what information

- do you lose? (limitation point of view)
- do you get rid of? (invariance point of view)

Which perspective is correct?



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Dr. Mubarak Shah

Re-think of histogram's application

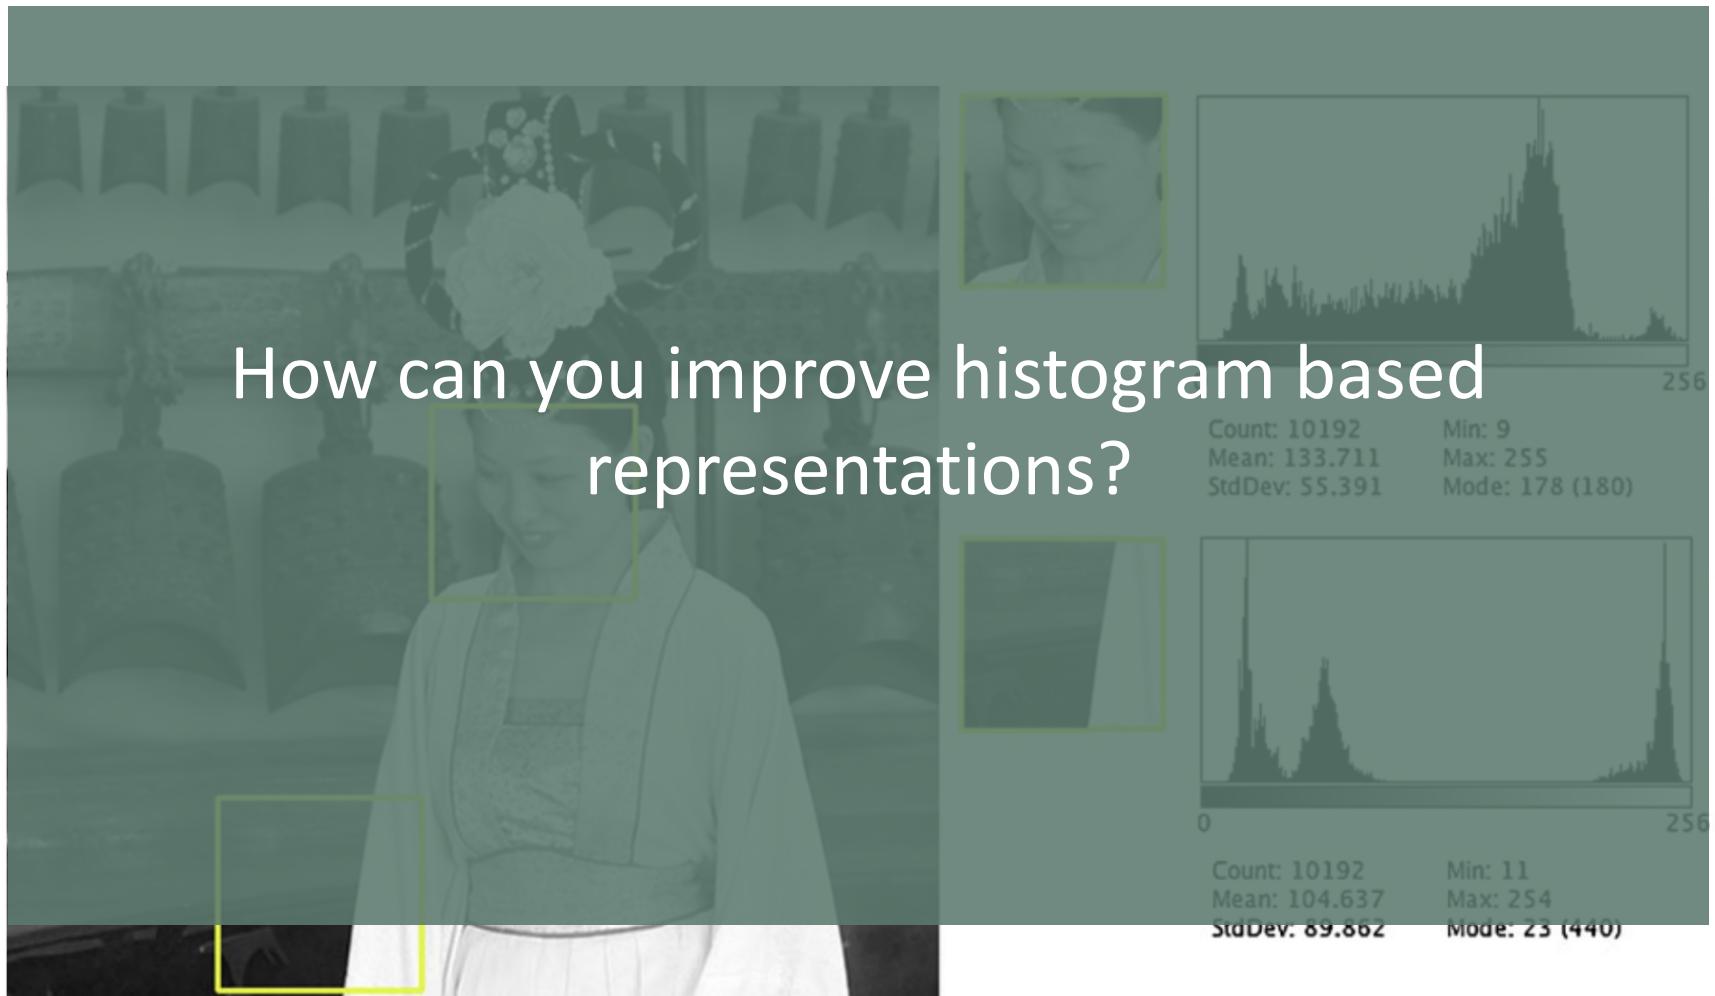
What is the direct limitation/invariance of a histogram? ie. what information

- do you lose? (limitation point of view)
- do you get rid of? (invariance point of view)

Both points of views can be valid, depending on the task! Examples:

- detection of camera orientation
- detection of cars in satellite imagery

Re-think of histogram's application



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Slide credit: Dr. Mubarak Shah

What we will learn today?

- Image sampling and quantization
- Image histograms
- **Images as functions**
- Linear systems (filters)
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

Images as discrete functions

- Images are usually **digital (discrete)**:
 - Sample the 2D space on a regular grid
- Represented as a matrix of integer values

pixel

62	79	23	119	120	05	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Images as coordinates

Cartesian coordinates

$$f[n, m] = \begin{bmatrix} \dots & & & \vdots \\ & f[-1, -1] & f[0, -1] & f[1, -1] \\ \dots & f[-1, 0] & \underline{f[0, 0]} & f[1, 0] & \dots \\ & f[-1, 1] & f[0, 1] & f[1, 1] & \ddots \\ \vdots & & & & \ddots \end{bmatrix}$$

Notation for discrete functions

A blue arrow points from the text "Notation for discrete functions" up towards the matrix, indicating that the notation $f[n, m]$ is used to represent discrete functions.

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Images as functions

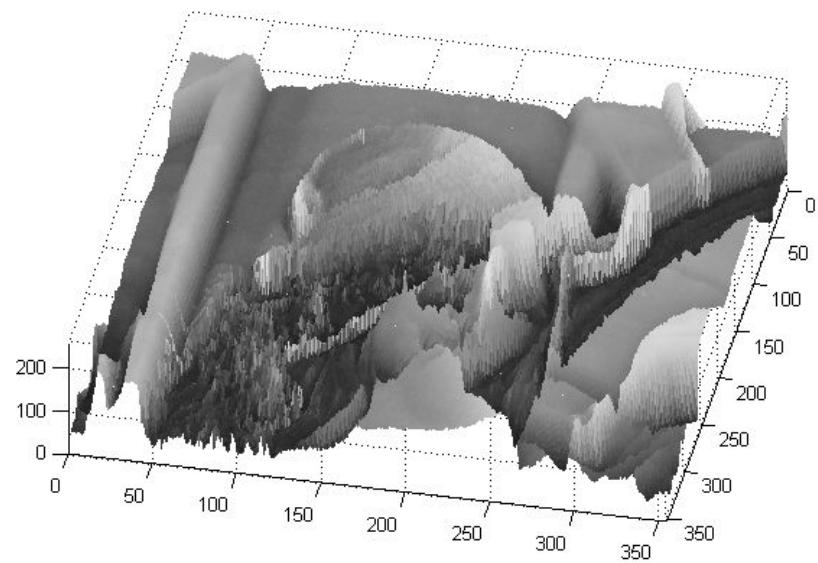
- An Image as a function f from \mathbb{R}^2 to \mathbb{R}^M :
 - $f(x, y)$ gives the **intensity** at position (x, y)
 - Defined over a rectangle, with a finite range:

$$f: [a,b] \times [c,d] \rightarrow [0,255]$$

 Domain
support

range

hna



Adapted from slides

Images as functions

- An Image as a function f from \mathbb{R}^2 to \mathbb{R}^M :
 - $f(x, y)$ gives the **intensity** at position (x, y)
 - Defined over a rectangle, with a finite range:

$$f: [a,b] \times [c,d] \rightarrow [0,255]$$

Domain
support

- A color image: $f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Histograms are a type of image function

- Histogram can be seen as another $f: \text{region} \rightarrow \text{vector mapping}$.
- In recognition tasks for computer vision, *image representation* typically plays a critical role.
- Throughout the course, we'll progressively discuss **more powerful representation** extraction techniques.

Histograms are a type of image function

- Histogram can be seen as another $f: \text{region} \rightarrow \text{vector mapping}$.
- In recognition tasks for computer vision, *image representation* typically plays a critical role.
- Throughout the course, we'll progressively discuss **more powerful representation** extraction techniques.

More *powerful* is not necessarily more *detailed*. (Think of the trade-offs in histograms, as a good example.)

Utility of a representation is typically task specific and plays a central role in the success of computer vision approaches.

What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- **Linear systems (filters)**
- Convolution and correlation

Some background reading:

Forsyth and Ponce, Computer Vision, Chapter 7

Systems and Filters

Filtering:

- Forming a new image whose pixel values are transformed from original pixel values

Goals:

- Goal is to extract useful information from images, or transform images into another domain where we can modify/enhance image properties
 - Features (edges, corners, blobs...)
 - super-resolution; in-painting; de-noising

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

System and Filters

- We define a system as a unit that converts an input function $f[n,m]$ into an output (or response) function $g[n,m]$, where (n,m) are the independent variables.
 - In the case for images, (n,m) represents the **spatial position in the image**.

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$

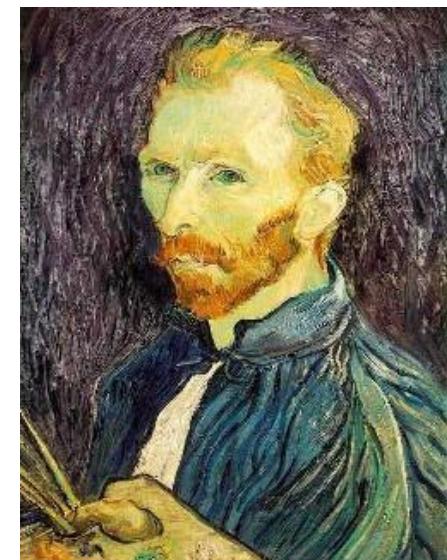
De-noising



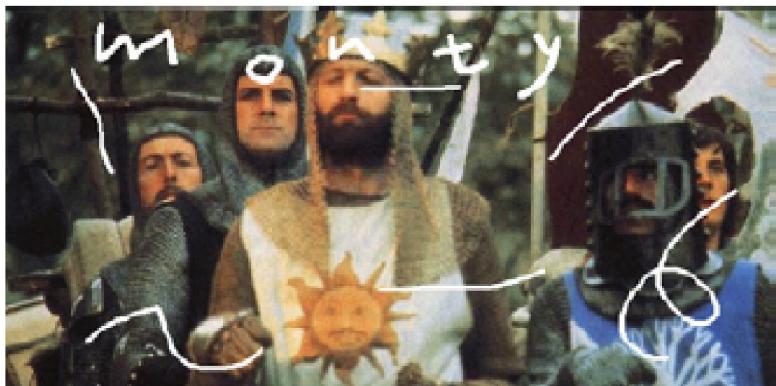
Salt and pepper noise



Super-resolution



In-painting



Bertamio et al

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

2D discrete-space systems (filters)

A *digital filter* is a system that performs mathematical operations on a sampled, discrete-time signal. [Wikipedia]

- S is the **system operator**, defined as a mapping or assignment of a member of the set of possible outputs $g[n,m]$ to each member of the set of possible inputs $f[n,m]$.

$$f[n, m] \rightarrow \boxed{\text{System } S} \rightarrow g[n, m]$$

$$g = S[f], \quad g[n, m] = S\{f[n, m]\}$$

$$f[n, m] \xrightarrow{S} g[n, m]$$

Filter example #1: Moving Average

2D discrete space moving average over a 3×3 window of neighborhood

$$\begin{aligned} g[n, m] &= \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l] \\ &= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l] \end{aligned}$$

$$h = \begin{matrix} & & h \\ \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} & \end{matrix}$$

Filter example #1: Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Filter example #1: Moving Average

$$F[x, y]$$

$$G[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

0	10									

$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Filter example #1: Moving Average

$$F[x, y]$$

$$G[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

0	10	20							

$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Filter example #1: Moving Average

$$F[x, y]$$

$$G[x, y]$$

$$(f * h)[m, n] = \sum f[k, l] h[m - k, n - l]$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna 

Filter example #1: Moving Average

$$F[x, y]$$

$$G[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Filter example #1: Moving Average

$$F[x, y]$$

$$G[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$(f * h)[m, n] = \sum_{k,l} f[k, l] h[m - k, n - l]$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Source: S. Seitz

Filter example #1: Moving Average

In summary:

- This filter “Replaces” each pixel with an average of its neighborhood.
- Achieve smoothing effect (remove sharp features)

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} h[\cdot, \cdot]$$

Filter example #1: Moving Average



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Filter example #2: Thresholding

- Image segmentation based on a simple threshold:

$$g[n, m] = \begin{cases} 255, & f[n, m] > 100 \\ 0, & \text{otherwise.} \end{cases}$$



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Properties of linear systems

- Amplitude properties:
 - Additivity

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

Properties of linear systems

- Amplitude properties:

- Additivity

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

- Homogeneity

$$S[\alpha f_i[n, m]] = \alpha S[f_i[n, m]]$$

Properties of linear systems

- Amplitude properties:

- Additivity

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

- Homogeneity

$$S[\alpha f_i[n, m]] = \alpha S[f_i[n, m]]$$

- Superposition

$$S[\alpha f_i[n, m] + \beta f_j[n, m]] = \alpha S[f_i[n, m]] + \beta S[f_j[n, m]]$$

Properties of linear systems

- Amplitude properties:

- Additivity

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

- Homogeneity

$$S[\alpha f_i[n, m]] = \alpha S[f_i[n, m]]$$

- Superposition

$$S[\alpha f_i[n, m] + \beta f_j[n, m]] = \alpha S[f_i[n, m]] + \beta S[f_j[n, m]]$$

- Stability (bounded-input, bounded-output (BIBO) stability)

$$|f[n, m]| \leq k \implies |g[n, m]| \leq ck$$

Properties of linear systems

- Amplitude properties:

- Additivity [Required]

$$S[f_i[n, m] + f_j[n, m]] = S[f_i[n, m]] + S[f_j[n, m]]$$

- Homogeneity [Required]

$$S[\alpha f_i[n, m]] = \alpha S[f_i[n, m]]$$

- Superposition (Generalizes additivity + homogeneity)

$$S[\alpha f_i[n, m] + \beta f_j[n, m]] = \alpha S[f_i[n, m]] + \beta S[f_j[n, m]]$$

- Stability [extra]

$$|f[n, m]| \leq k \implies |g[n, m]| \leq ck$$

- Invertibility [extra]

$$S^{-1}[S[f_i[n, m]]] = f_i[n, m]$$

Properties of linear systems

- Spatial properties
 - Causality [extra]

for $n < n_0, m < m_0$, if $f[n, m] = 0 \implies g[n, m] = 0$

- Shift invariance [extra]

Idea = same input at a *later time* yields the same response

$$f[n - n_0, m - m_0] \xrightarrow{\mathcal{S}} g[n - n_0, m - m_0]$$

Is the moving average system is shift invariant?

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$F[x, y]$

$G[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Is the moving average system is shift invariant?

$$f[n, m] \xrightarrow{s} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$$f[n - n_0, m - m_0]$$

$$\xrightarrow{s} \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[(n - n_0) - k, (m - m_0) - l]$$

$$= g[n - n_0, m - m_0]$$

Yes!

Is the moving average system causal?

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

$F[x, y]$

$G[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

for $n < n_0, m < m_0$, if $f[n, m] = 0 \implies g[n, m] = 0$

Adapted from slide

Is the moving average system causal?

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m] = \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

NO! It *looks*

into future

$[x, y]$

$G[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

for $n < n_0, m < m_0$, if $f[n, m] = 0 \implies g[n, m] = 0$

Adapted from slide

Linear Systems (filters)

$$f[n, m] \rightarrow \boxed{\text{System } S} \rightarrow g[n, m]$$

- Linear filtering:
 - Form a new image whose pixels are a weighted sum of original pixel values
 - Use the same set of weights at each point
- **S** is a linear system (function) iff it *S satisfies*

$$S[\alpha f_i[n, m] + \beta f_j[h, m]] = \alpha S[f_i[n, m]] + \beta S[f_j[h, m]]$$

superposition property

Linear Systems (filters)

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$

- Is the moving average a linear system?

(Exercise)

- Is thresholding a linear system?

- $f_1[n, m] + f_2[n, m] > T$
- $f_1[n, m] < T$
- $f_2[n, m] < T$

No!

A more interesting exercise

$$y[n] = n \cdot x[n]$$

- linear?
- shift/time-invariant?
- causal?
- stable?

An analysis can be found here:

<https://www.youtube.com/watch?v=fO6RRXDacUw>

What we will learn today?

- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:
Forsyth and Ponce, Computer Vision, Chapter 7

1D Discrete convolution (symbol: $*$)

We are going to convolve a **function f** with a **filter h** .

$$g[n] = \sum_k f[k]h[n - k]$$



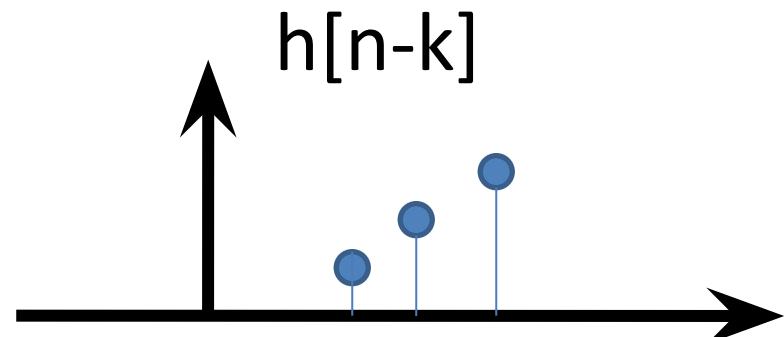
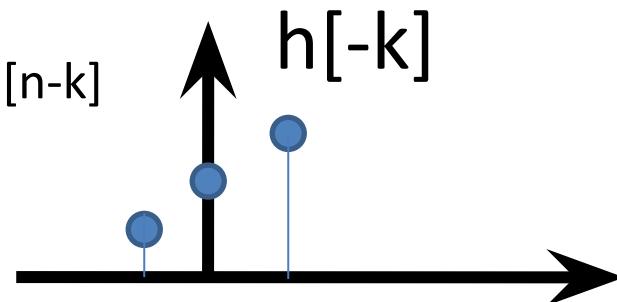
Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

1D Discrete convolution (symbol \ast)

We are going to convolve a **function f** with a **filter h**.

$$g[n] = \sum_k f[k]h[n - k]$$

We first need to calculate the shifted filter $h[n-k]$

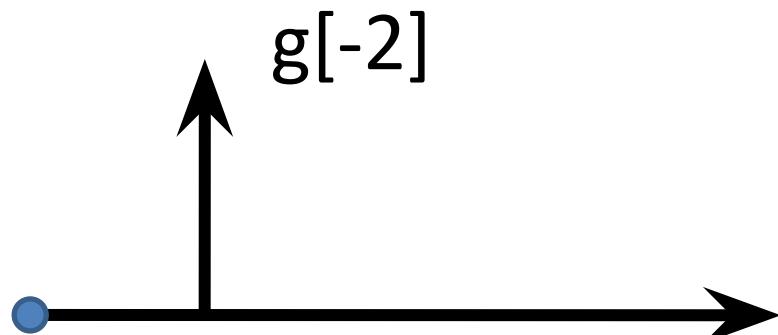
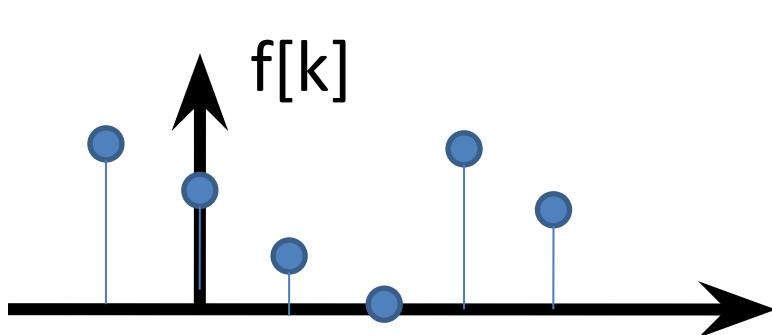
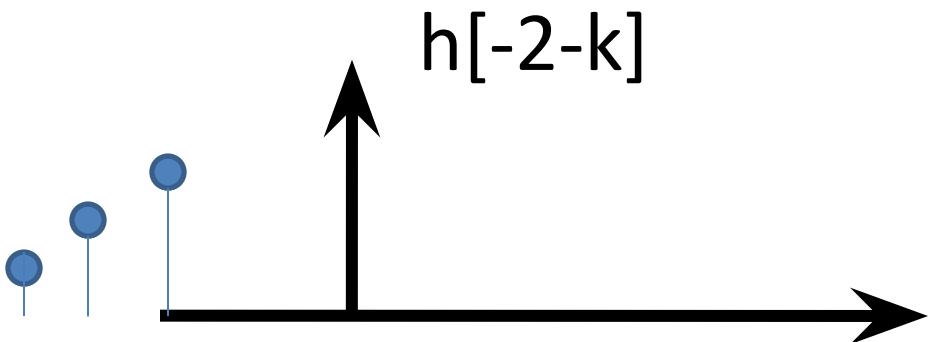


Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Discrete convolution (symbol: $*$)

We are going to convolve a function f with a filter h .

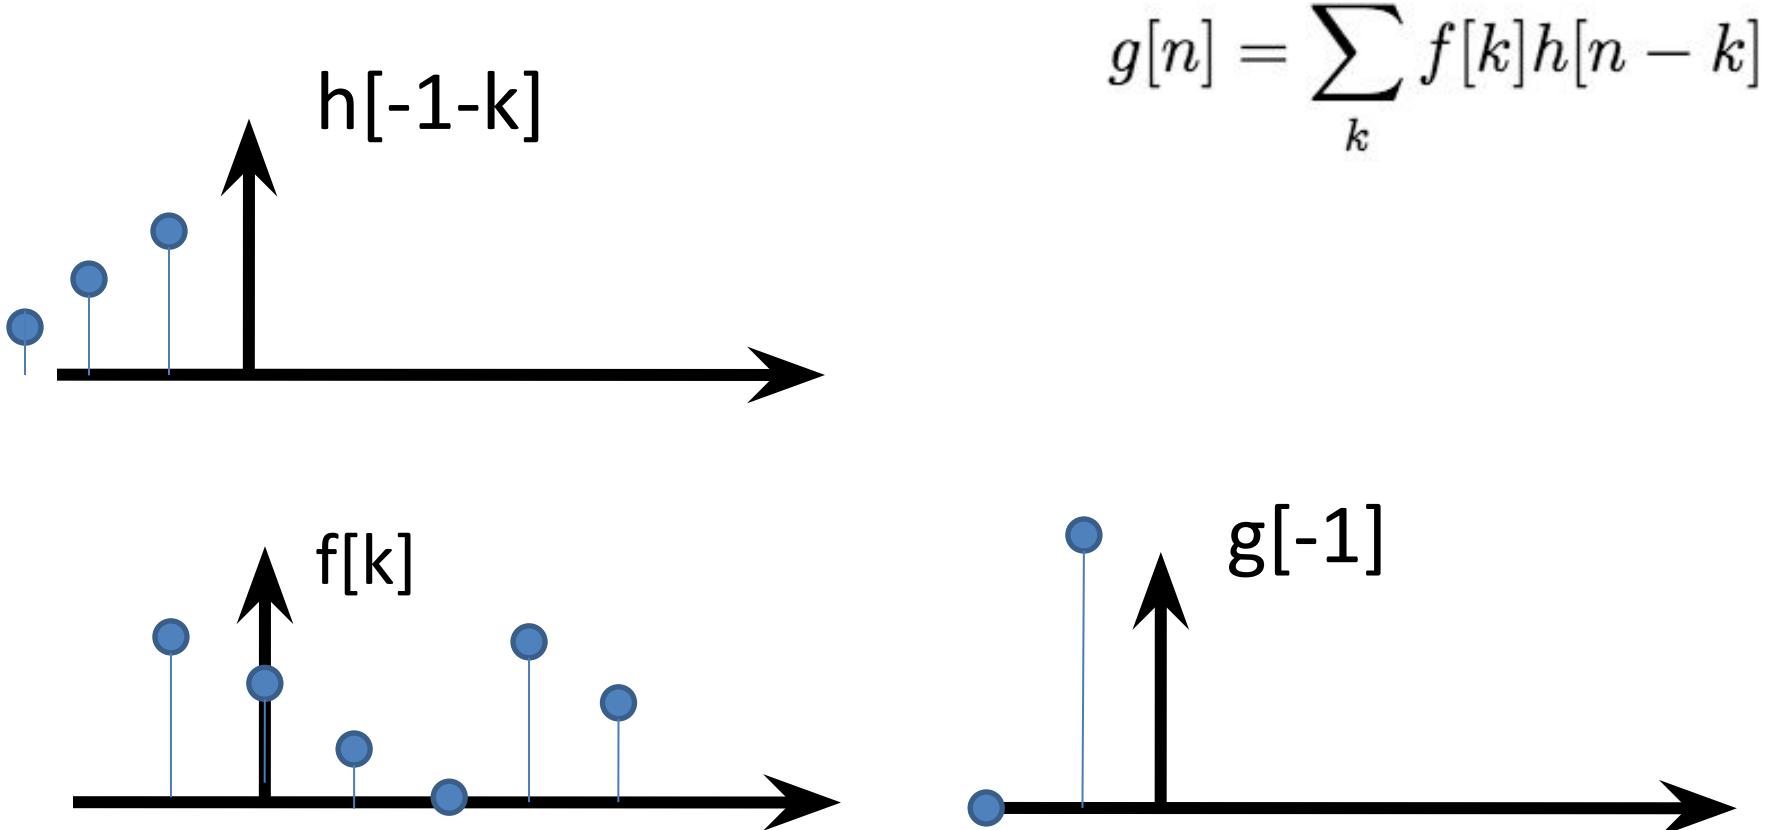
$$g[n] = \sum_k f[k]h[n - k]$$



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Discrete convolution (symbol: $*$)

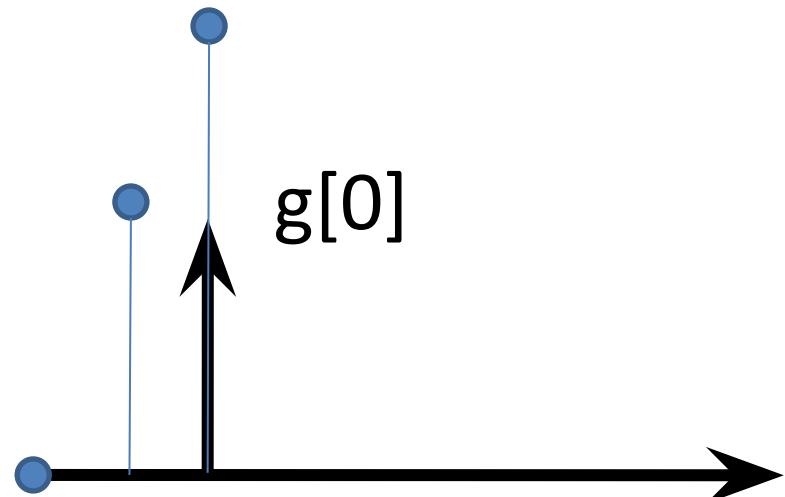
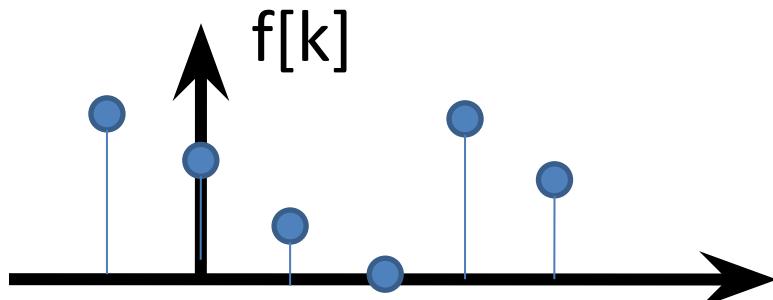
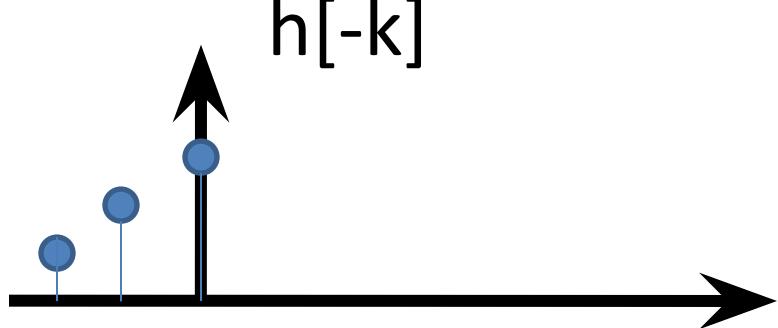
We are going to convolve a function f with a filter h .



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Discrete convolution (symbol: $*$)

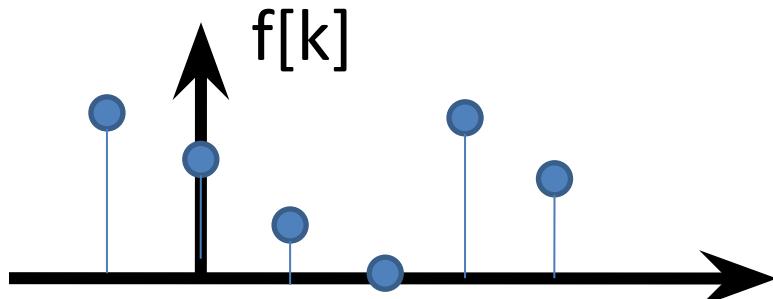
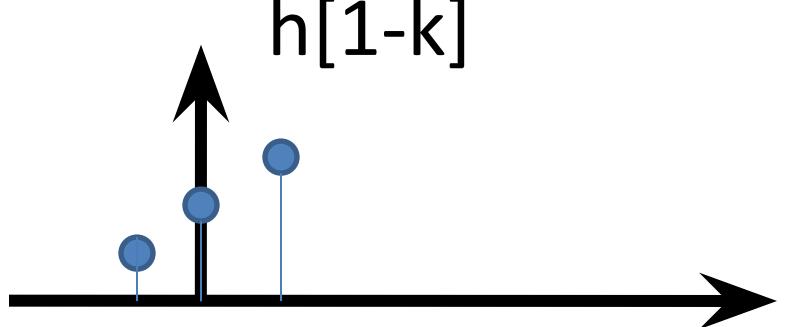
We are going to convolve a function f with a filter h .



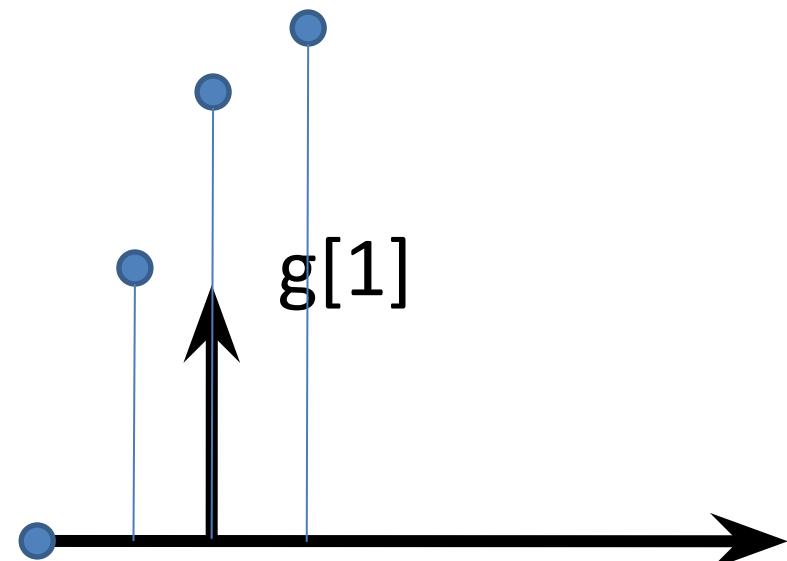
$$g[n] = \sum_k f[k]h[n - k]$$

Discrete convolution (symbol: $*$)

We are going to convolve a function f with a filter h .



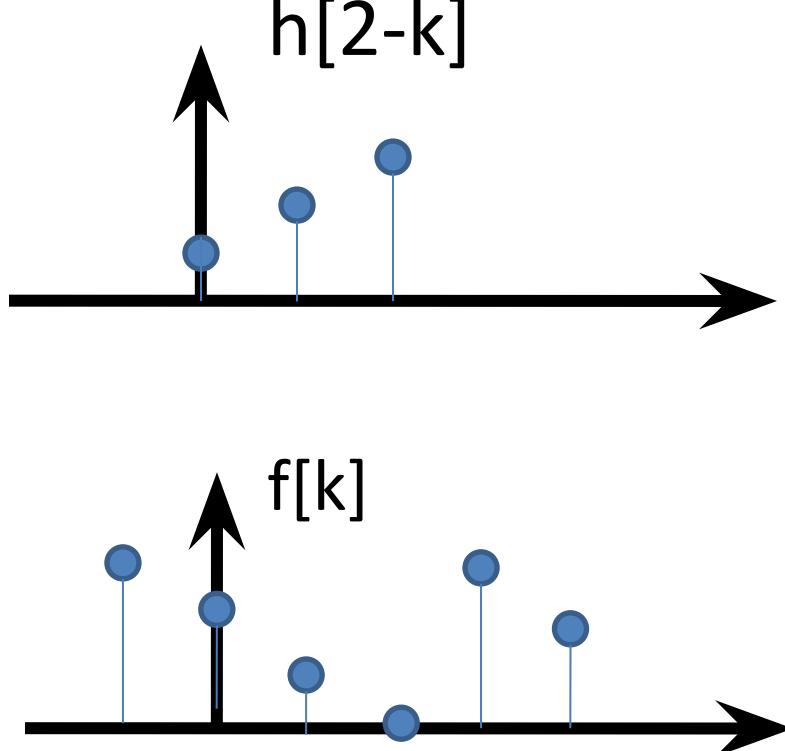
$$g[n] = \sum_k f[k]h[n - k]$$



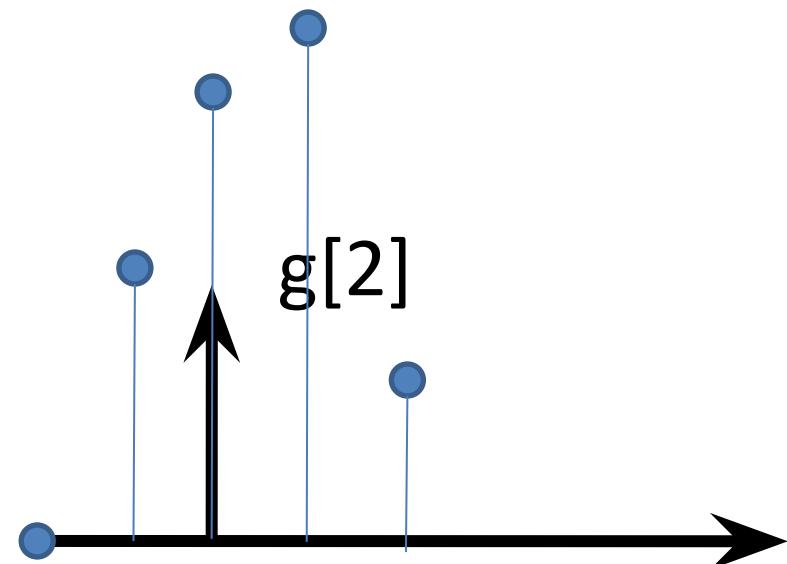
Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Discrete convolution (symbol: $*$)

We are going to convolve a function f with a filter h .



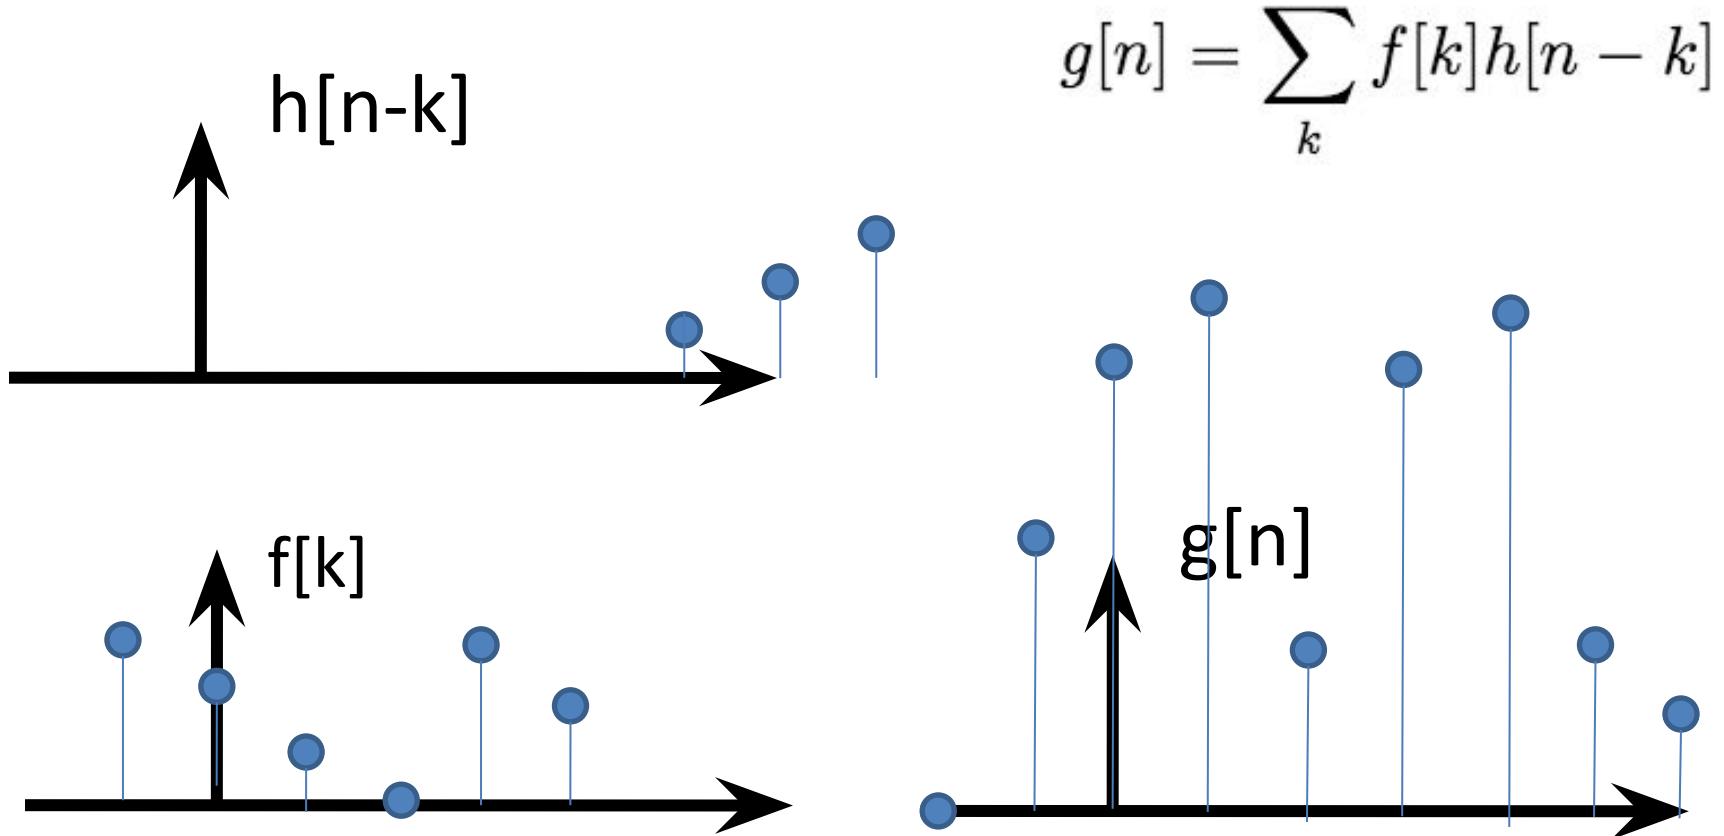
$$g[n] = \sum_k f[k]h[n - k]$$



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Discrete convolution (symbol: $*$)

We are going to convolve a function f with a filter h .



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Discrete convolution (symbol: $*$)

In summary, the steps for discrete convolution are:

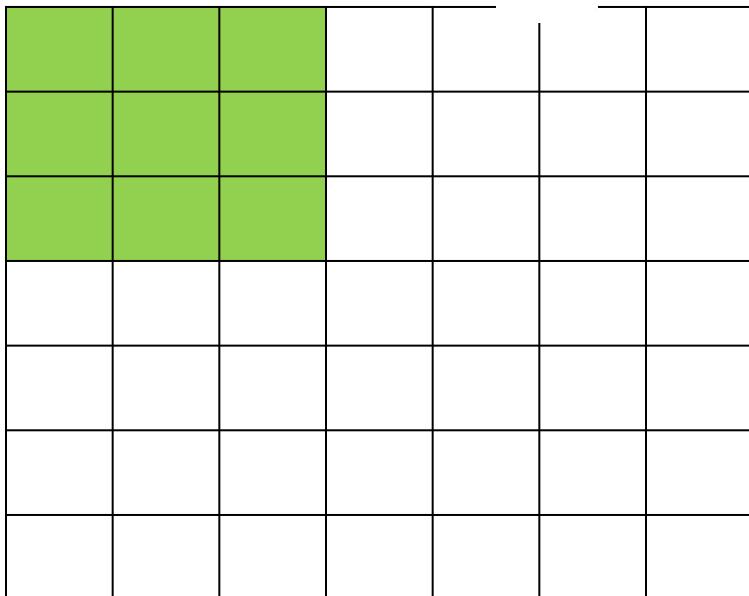
- Fold $h[k,l]$ about origin to form $h[-k]$
- Shift the folded results by n to form $h[n - k]$
- Multiply $h[n - k]$ by $f[k]$
- Sum over all k
- Repeat for every n

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.
(flip the kernel in 2 axis.)

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



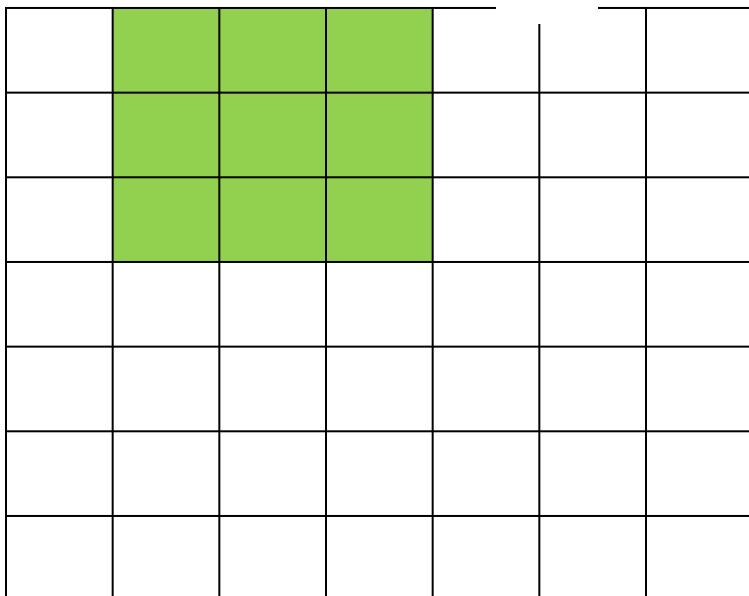
Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.
(flip the kernel in 2 axis.)

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



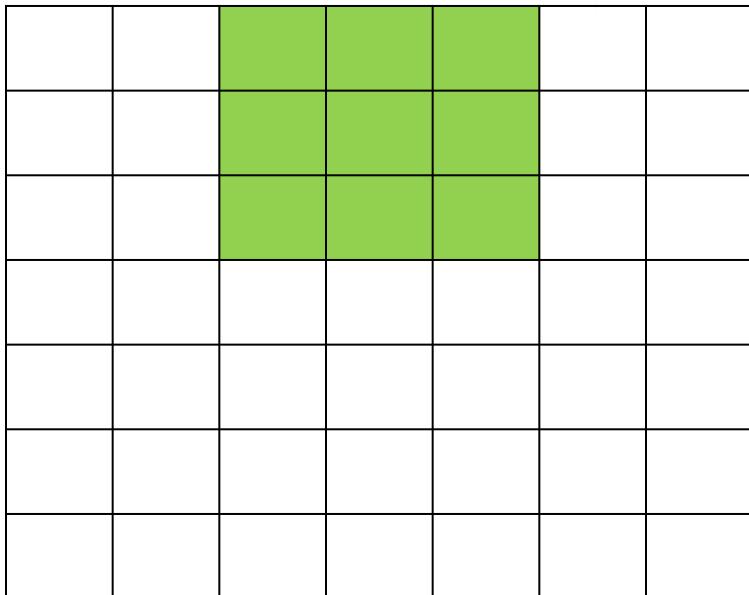
Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.
(flip the kernel in 2 axis.)

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



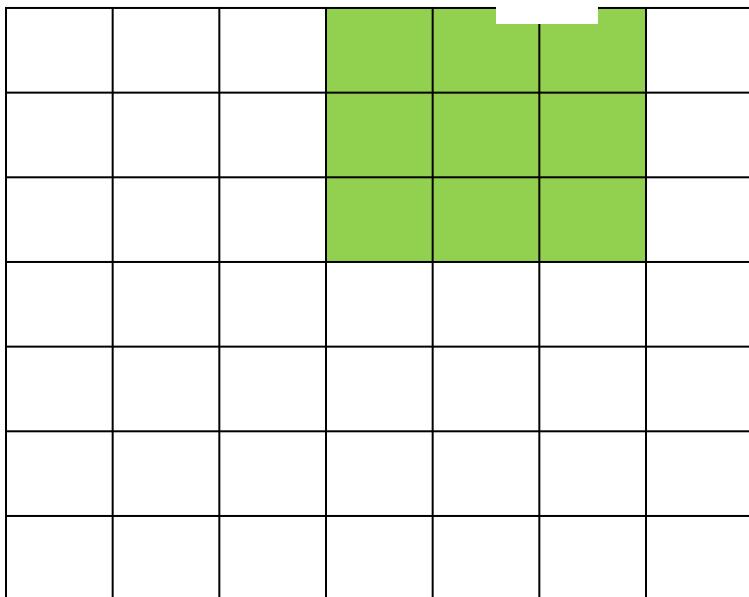
Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.
(flip the kernel in 2 axis.)

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



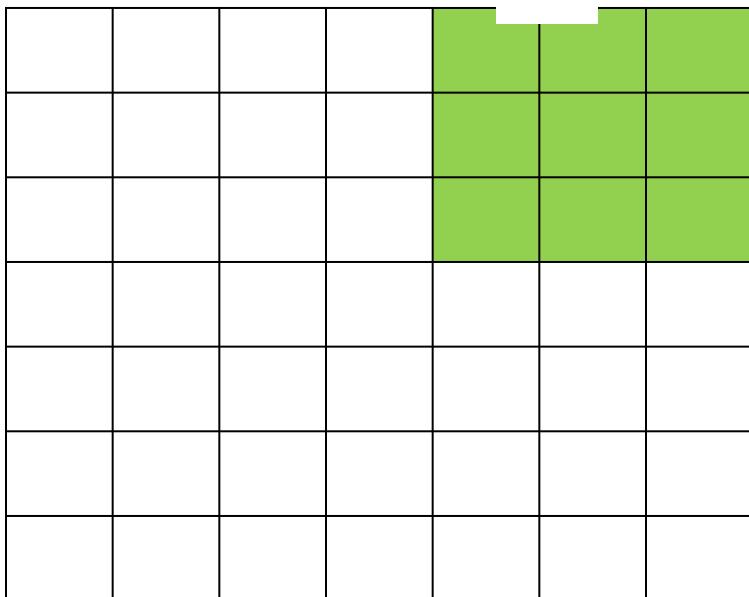
Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.
(flip the kernel in 2 axis.)

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



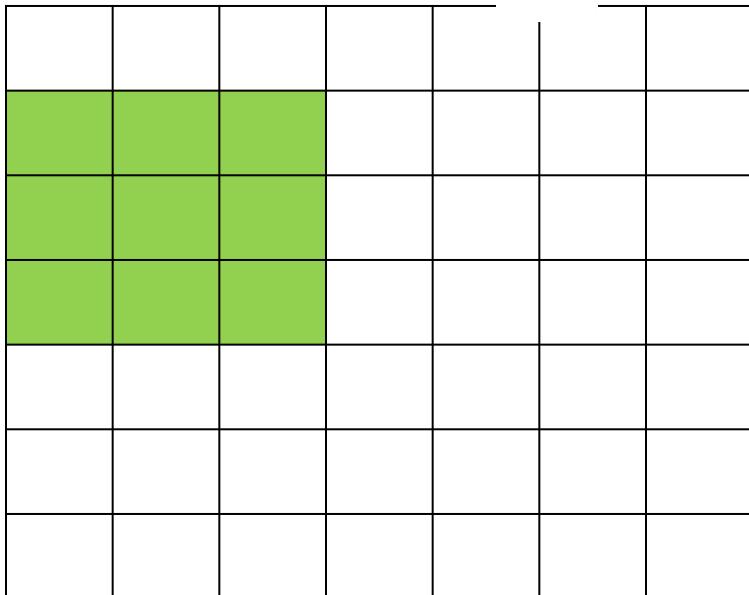
Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.

2D convolution

2D convolution is very similar to 1D.

- The main difference is that we now have to iterate over 2 axis instead of 1.
(flip the kernel in 2 axis.)

$$f[n, m] * h[n, m] = \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} f[k, l] h[n - k, m - l]$$



Assume we have a filter($h[,]$) that is 3x3. and an image ($f[,]$) that is 7x7.

2D convolution example

1	2	3
4	5	6
7	8	9

Input

m
n

-1	0	1
-1	-2	-1
0	0	0
1	1	1

Kernel

-13	-20	-17
-18	-24	-18
13	20	17

Output

2D convolution example

1	2	1		
0	0	0	2	3
-1	-2	-1	4	5
7	8	9		

$$\begin{aligned} &= x[-1,-1] \cdot h[1,1] + x[0,-1] \cdot h[0,1] + x[1,-1] \cdot h[-1,1] \\ &\quad + x[-1,0] \cdot h[1,0] + x[0,0] \cdot h[0,0] + x[1,0] \cdot h[-1,0] \\ &\quad + x[-1,1] \cdot h[1,-1] + x[0,1] \cdot h[0,-1] + x[1,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 0 \cdot 0 + 1 \cdot 0 + 2 \cdot 0 + 0 \cdot (-1) + 4 \cdot (-2) + 5 \cdot (-1) = -13 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn

2D convolution example

1	2	1
0	0	0
1	2	3
-1	-2	-1
4	5	6
7	8	9

$$\begin{aligned} &= x[0,-1] \cdot h[1,1] + x[1,-1] \cdot h[0,1] + x[2,-1] \cdot h[-1,1] \\ &\quad + x[0,0] \cdot h[1,0] + x[1,0] \cdot h[0,0] + x[2,0] \cdot h[-1,0] \\ &\quad + x[0,1] \cdot h[1,-1] + x[1,1] \cdot h[0,-1] + x[2,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 1 \cdot 0 + 2 \cdot 0 + 3 \cdot 0 + 4 \cdot (-1) + 5 \cdot (-2) + 6 \cdot (-1) = -20 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn

2D convolution example

	1	2	1
1	0	0	0
4	-1	-2	-1
7	8	9	

$$\begin{aligned} &= x[1,-1] \cdot h[1,1] + x[2,-1] \cdot h[0,1] + x[3,-1] \cdot h[-1,1] \\ &\quad + x[1,0] \cdot h[1,0] + x[2,0] \cdot h[0,0] + x[3,0] \cdot h[-1,0] \\ &\quad + x[1,1] \cdot h[1,-1] + x[2,1] \cdot h[0,-1] + x[3,1] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 2 \cdot 0 + 3 \cdot 0 + 0 \cdot 0 + 5 \cdot (-1) + 6 \cdot (-2) + 0 \cdot (-1) = -17 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn

2D convolution example

1	2	1	2	3
0	0	0	5	6
-1	-2	-1	8	9

$$\begin{aligned} &= x[-1,0] \cdot h[1,1] + x[0,0] \cdot h[0,1] + x[1,0] \cdot h[-1,1] \\ &\quad + x[-1,1] \cdot h[1,0] + x[0,1] \cdot h[0,0] + x[1,1] \cdot h[-1,0] \\ &\quad + x[-1,2] \cdot h[1,-1] + x[0,2] \cdot h[0,-1] + x[1,2] \cdot h[-1,-1] \\ &= 0 \cdot 1 + 1 \cdot 2 + 2 \cdot 1 + 0 \cdot 0 + 4 \cdot 0 + 5 \cdot 0 + 0 \cdot (-1) + 7 \cdot (-2) + 8 \cdot (-1) = -18 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn

2D convolution example

1	2	1
1	2	3
0	0	0
4	5	6
-1	-2	-1
7	8	9

$$\begin{aligned} &= x[0,0] \cdot h[1,1] + x[1,0] \cdot h[0,1] + x[2,0] \cdot h[-1,1] \\ &\quad + x[0,1] \cdot h[1,0] + x[1,1] \cdot h[0,0] + x[2,1] \cdot h[-1,0] \\ &\quad + x[0,2] \cdot h[1,-1] + x[1,2] \cdot h[0,-1] + x[2,2] \cdot h[-1,-1] \\ &= 1 \cdot 1 + 2 \cdot 2 + 3 \cdot 1 + 4 \cdot 0 + 5 \cdot 0 + 6 \cdot 0 + 7 \cdot (-1) + 8 \cdot (-2) + 9 \cdot (-1) = -24 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn

2D convolution example

1	1	2	3	1
4	0	5	6	0
7	-1	8	-2	9

$$\begin{aligned} &= x[1,0] \cdot h[1,1] + x[2,0] \cdot h[0,1] + x[3,0] \cdot h[-1,1] \\ &\quad + x[1,1] \cdot h[1,0] + x[2,1] \cdot h[0,0] + x[3,1] \cdot h[-1,0] \\ &\quad + x[1,2] \cdot h[1,-1] + x[2,2] \cdot h[0,-1] + x[3,2] \cdot h[-1,-1] \\ &= 2 \cdot 1 + 3 \cdot 2 + 0 \cdot 1 + 5 \cdot 0 + 6 \cdot 0 + 0 \cdot 0 + 8 \cdot (-1) + 9 \cdot (-2) + 0 \cdot (-1) = -18 \end{aligned}$$

-13	-20	-17
-18	-24	-18
13	20	17

Output

Slide credit: Song Ho Ahn

Convolution in 2D - examples



$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 1 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} \quad = \quad ?$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution in 2D - examples



*

$$\begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 1 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array}$$

=



Original

Filtered
(no change)

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

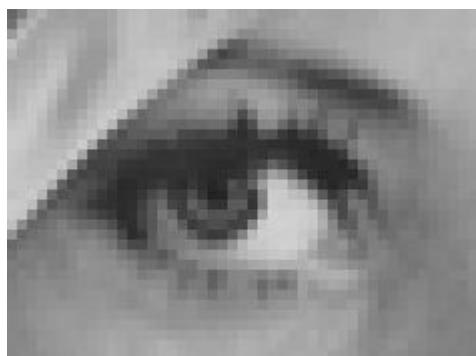
Convolution in 2D - examples



$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 1 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} \quad = \quad ?$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution in 2D - examples

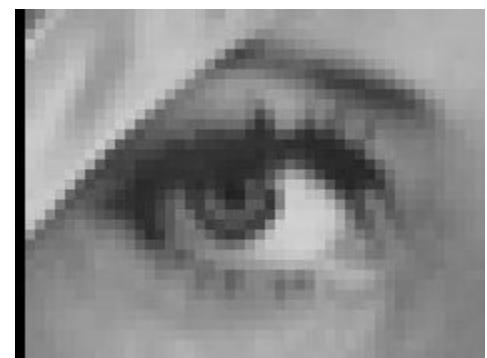


Original

*

•0	•0	•0
•0	•0	•1
•0	•0	•0

=



Shifted right
By 1 pixel

•0	•0	•0
•1	•0	•0
•0	•0	•0

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution in 2D - examples



Original

$$\text{Original} * \begin{matrix} 1 & & \\ & 1 & \\ & & 1 \end{matrix} = ?$$

A diagram illustrating a 2D convolution operation. On the left is a grayscale image of a face labeled "Original". To its right is a convolution kernel represented as a 3x3 matrix with ones in all positions. An asterisk (*) indicates the operation, followed by an equals sign (=) and a question mark (?), indicating the result of the convolution.

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution in 2D - examples



$$\text{Original} \quad * \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array} = \text{Blur (with a box filter)}$$



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution in 2D - examples



Original

$$\begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 2 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix}$$

-

$$\frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix}$$

= ?

(Note that filter sums to 1)

“details of the image”

$$\begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix}$$

+

$$\begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix}$$

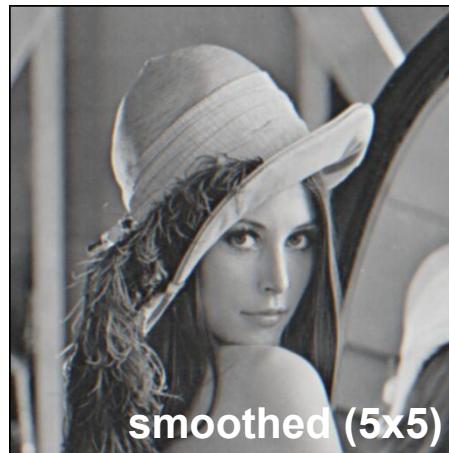
-

$$\frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix}$$

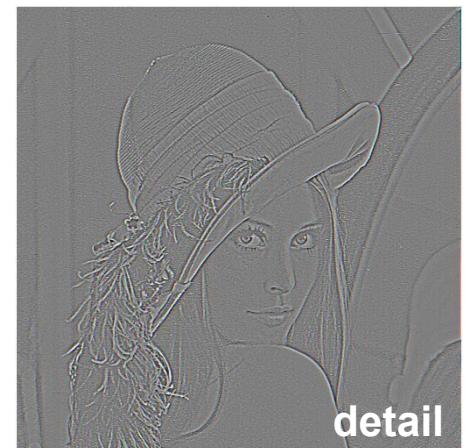
- What does blurring take away?



original



smoothed (5x5)



detail

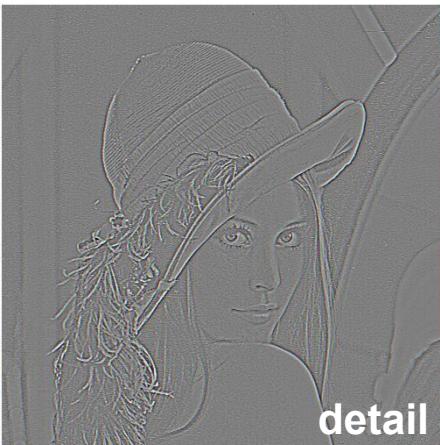
-

=



original

+ a



detail

=



sharpened

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution in 2D – Sharpening filter



Original

$$\begin{matrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 2 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{matrix}$$

-

$$\begin{matrix} 1 \\ 9 \mid \end{matrix} \begin{matrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{matrix}$$

=



Sharpening filter: Accentuates differences with local average

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Image support and edge effect

- A computer will only convolve **finite support signals**.
 - That is: images that are zero for n, m outside some rectangular region
- numpy's convolution performs 2D convolution of finite-support signals.

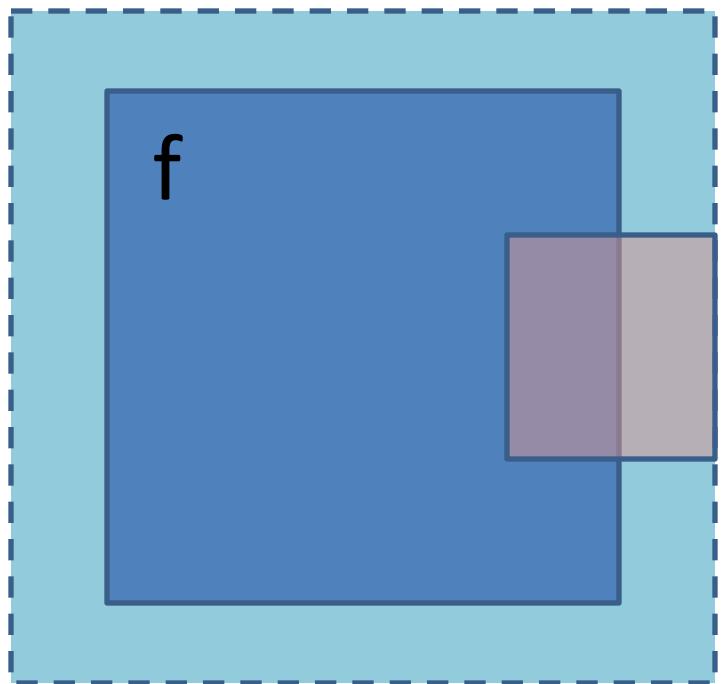
The diagram illustrates the convolution of two finite-support signals. On the left, a large blue square represents the input image, labeled $N_1 \times M_1$. In the center, a smaller red square represents the kernel, labeled $N_2 \times M_2$. To the right of the kernel is an equals sign. To the right of the equals sign is the result of the convolution, which is a green square divided into two horizontal bands: a top band of width $M_2 - 1$ and a bottom band of width $M_1 - M_2 + 1$. The total width of the result is $(M_1 - M_2 + 1) \times (N_1 - N_2 + 1)$.

$$\text{Input: } N_1 \times M_1$$
$$\text{Kernel: } N_2 \times M_2$$
$$\text{Result: } (N_1 - N_2 + 1) \times (M_1 - M_2 + 1)$$

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Image support and edge effect

- A computer will only convolve **finite support signals**.
- What happens at the edge?



- zero “padding”
 - edge value replication
 - flip extension
 - more (beyond the scope of this class)
- > Matlab conv2 uses
zero-padding

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

What we will learn today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Some background reading:
Forsyth and Ponce, Computer Vision, Chapter 7

(Cross) correlation (symbol: $\ast\ast$)

Cross correlation of two 2D signals $f[n,m]$ and $g[n,m]$

$$r_{fg}[k, l] \triangleq \sum_{n=-\infty}^{\infty} \sum_{m=-\infty}^{\infty} f[n, m] g^*[n - k, m - l]$$

$$= \sum_{n=-\infty}^{\infty} \sum_{m=-\infty}^{\infty} f[n + k, m + l] g^*[n, m], \quad k, l \in \mathbb{Z}.$$

(k, l) is called the **lag**

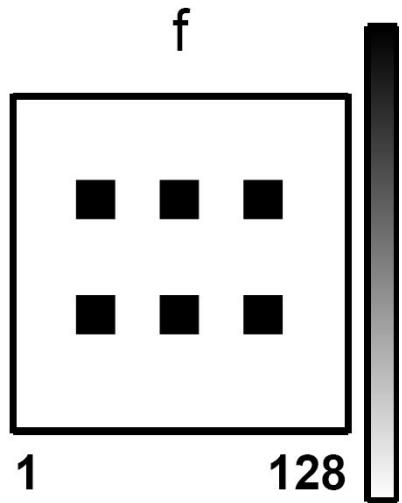
- Equivalent to a convolution without the flip

$$r_{fg}[n, m] = f[n, m] * g^*[-n, -m]$$

(g^* is defined as the *complex conjugate* of g . In this class, $g(n,m)$ are real numbers, hence $g^*=g$.)

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

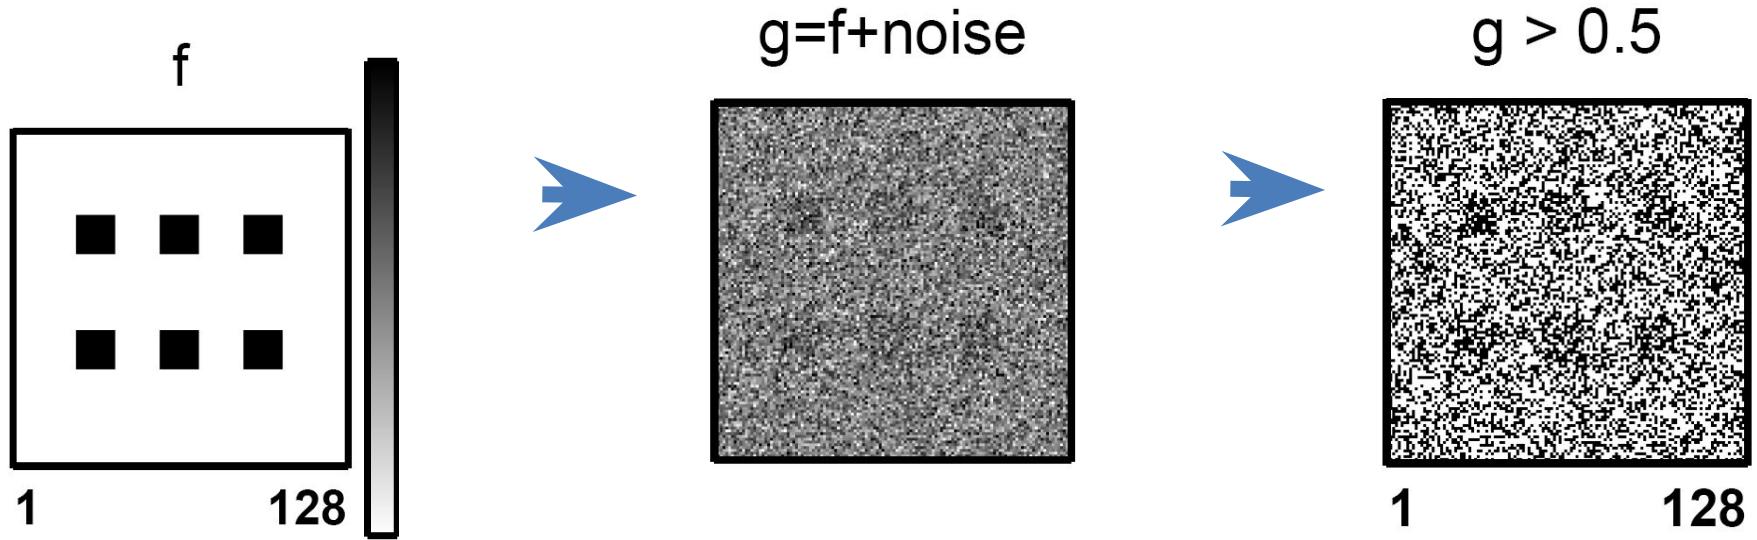
(Cross) correlation – example



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

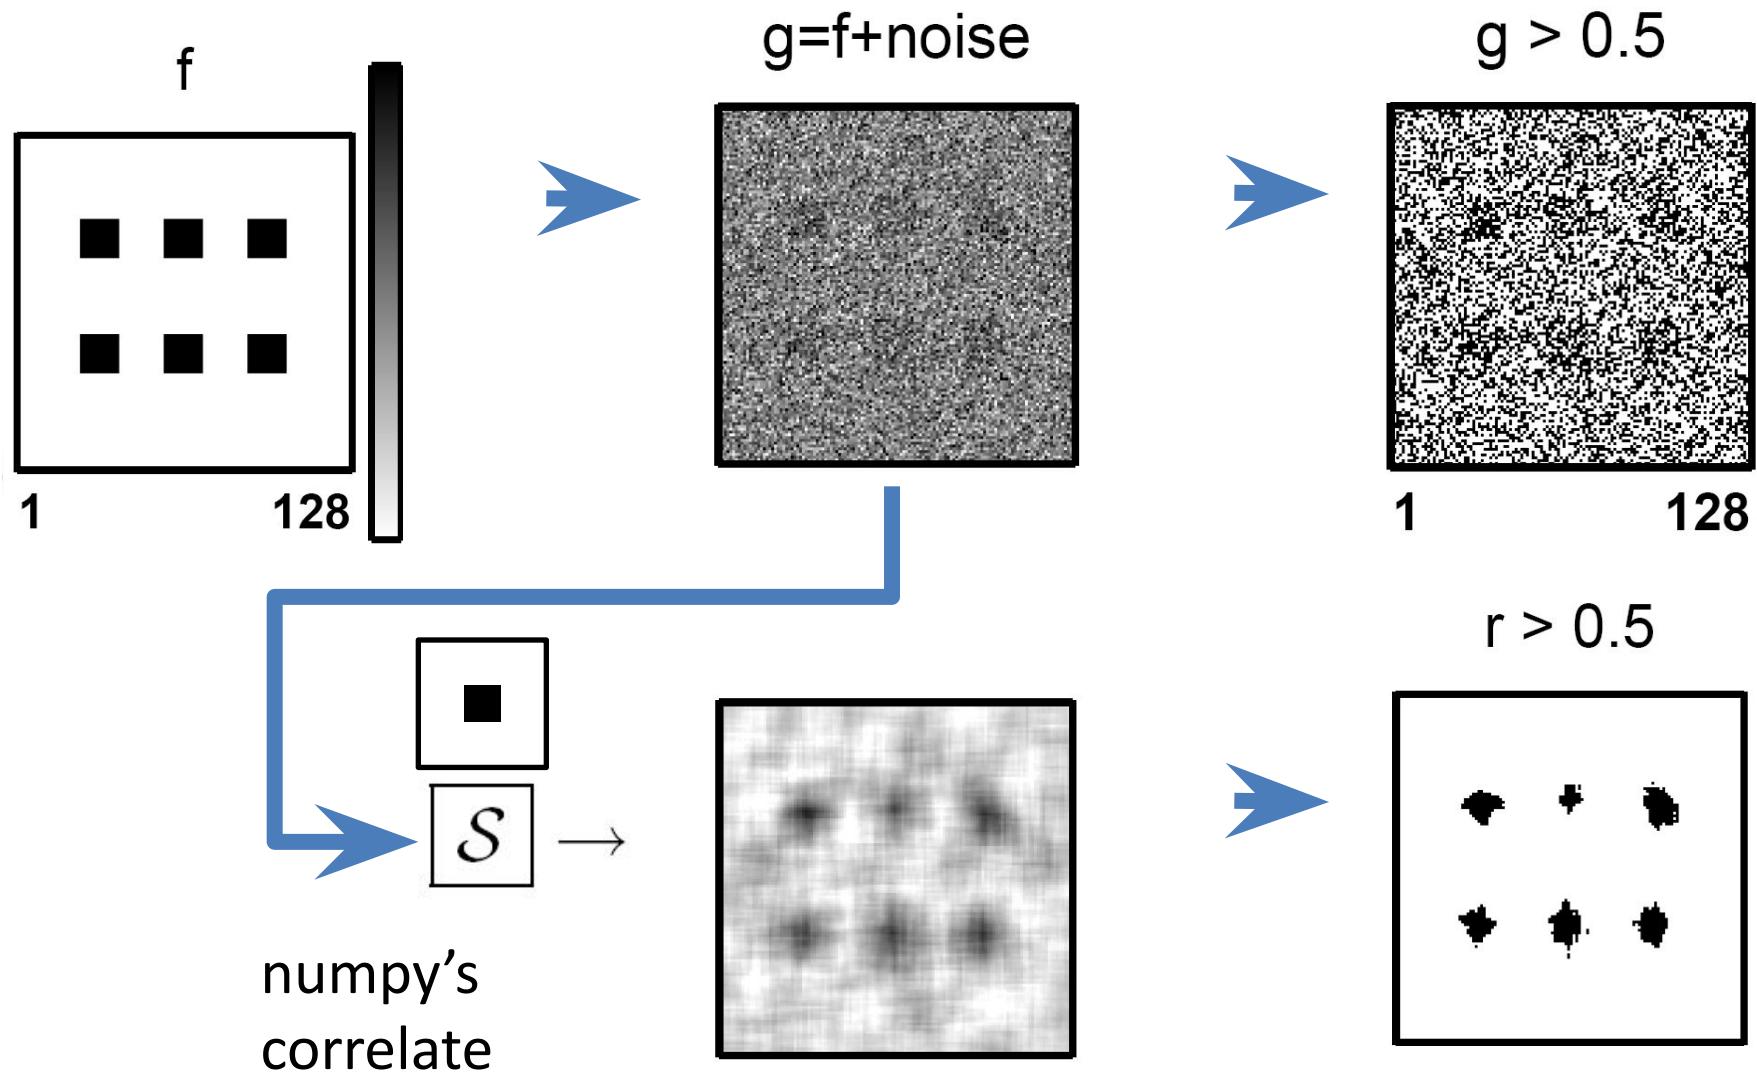
Courtesy of J. Fessler

(Cross) correlation – example



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

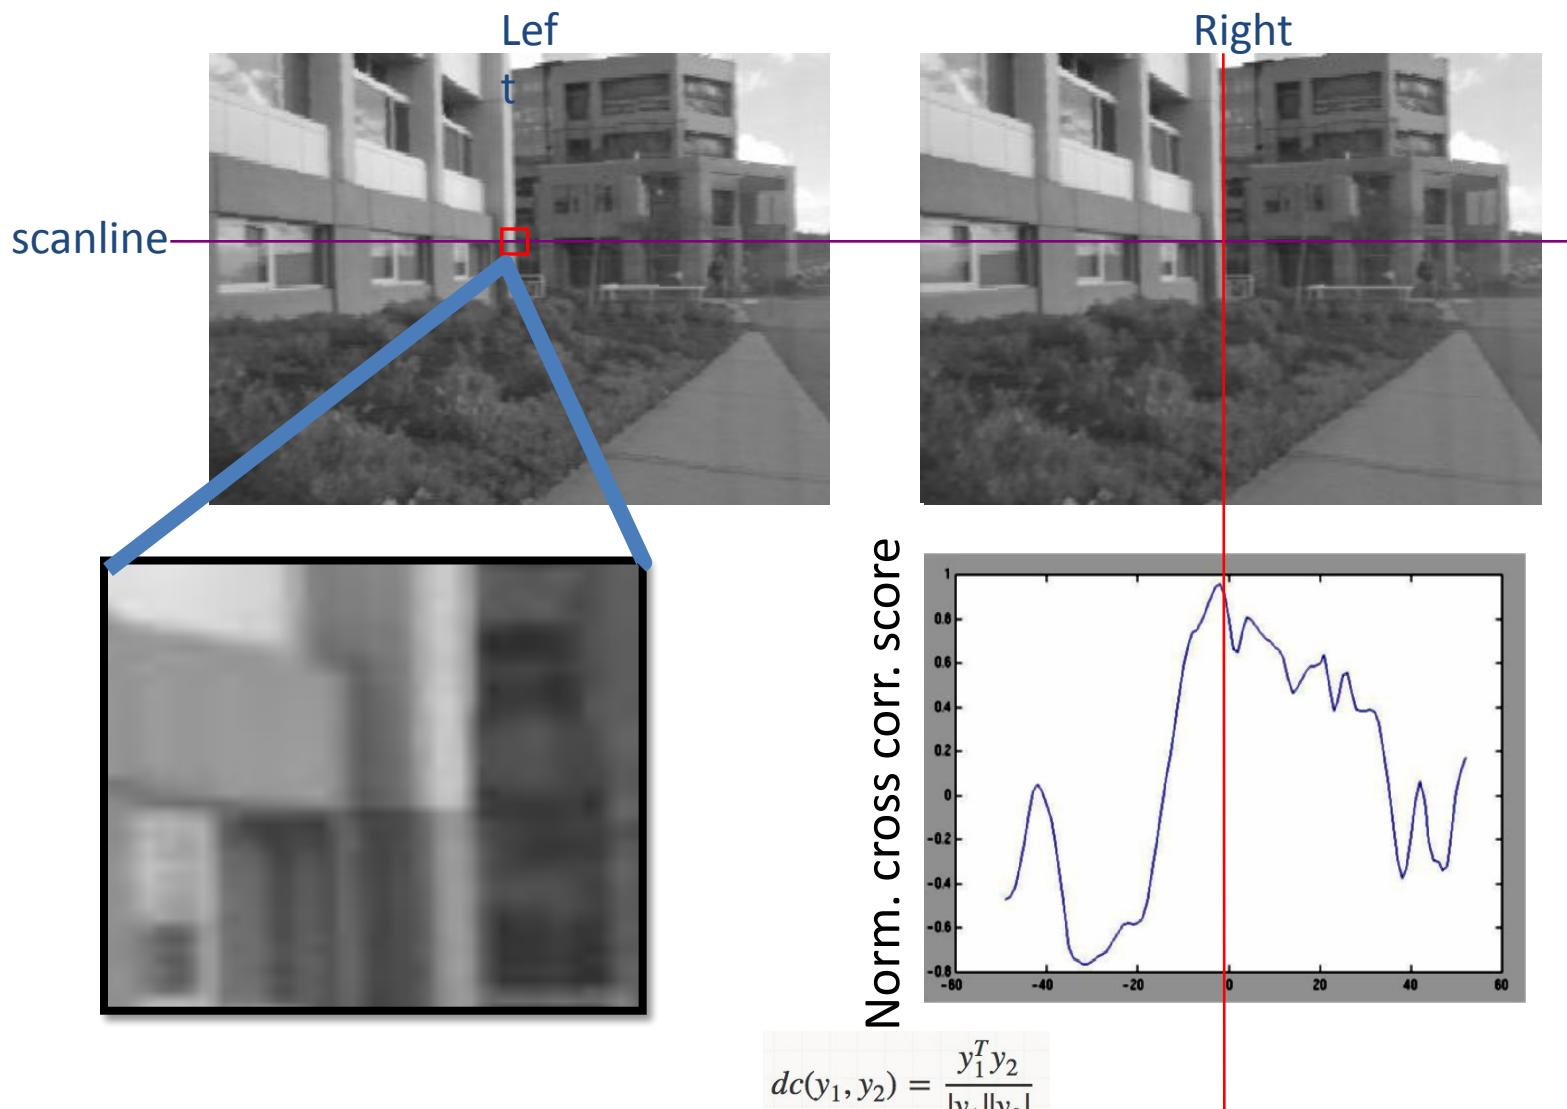
(Cross) correlation – example



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

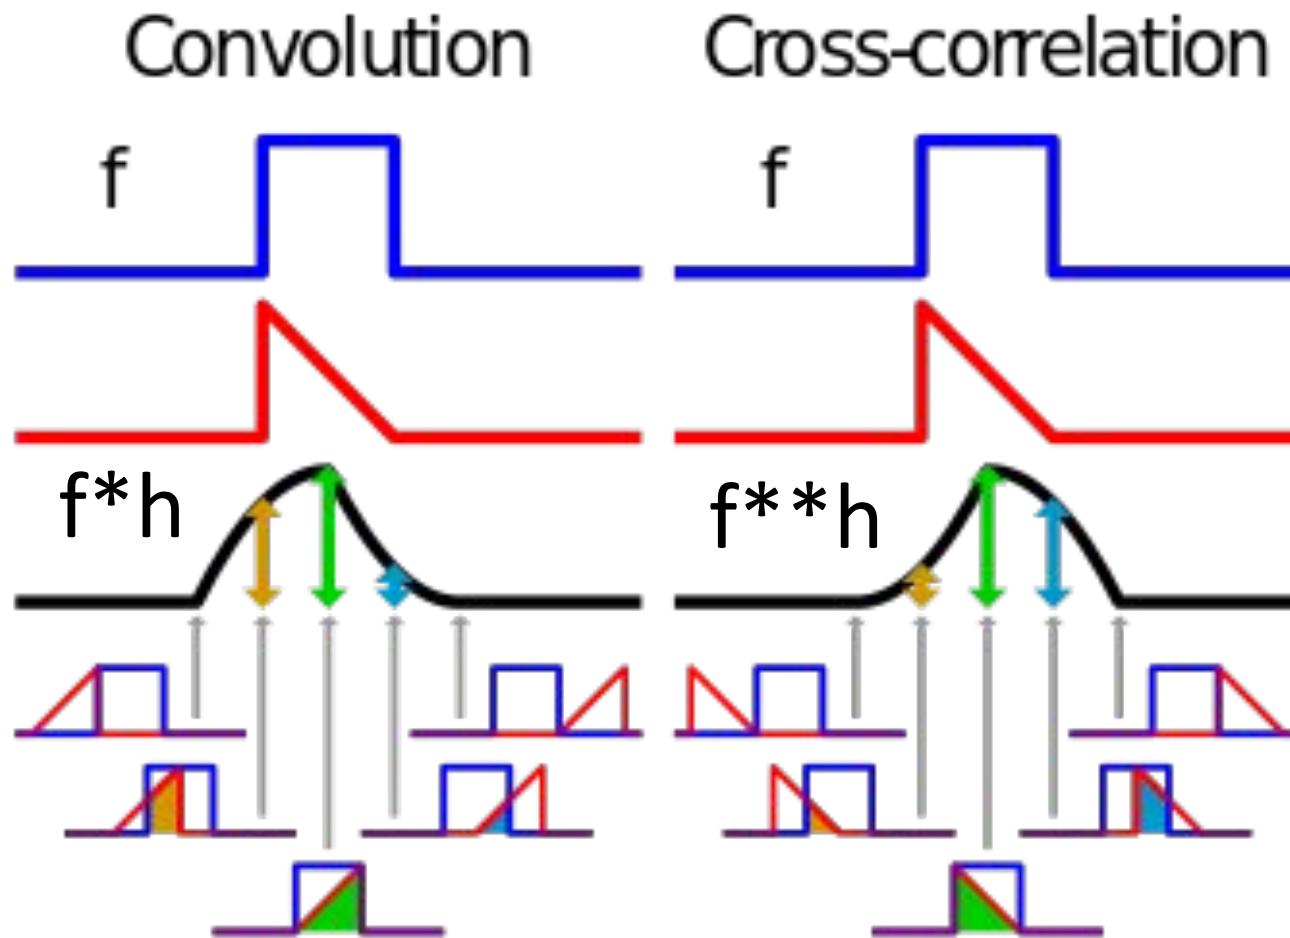
Courtesy of J. Fessler

(Cross) correlation – example



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution vs. (Cross) Correlation



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna



Cross Correlation Application: Vision system for TV remote control

- uses template matching

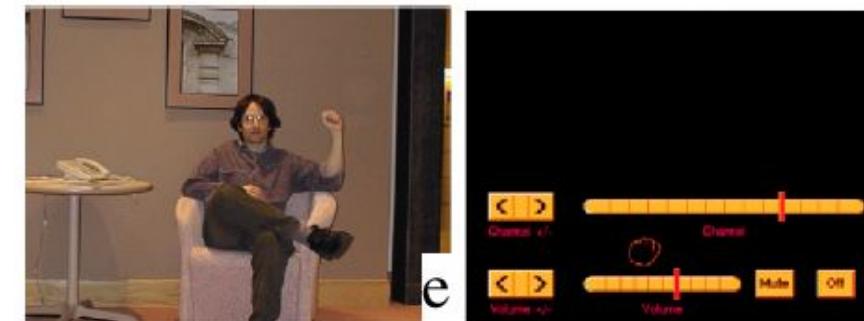
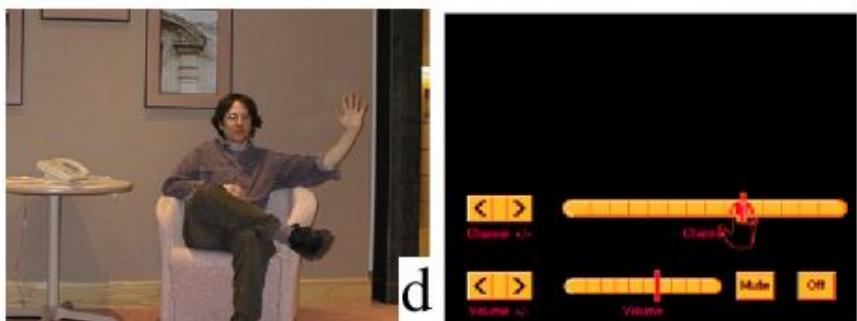
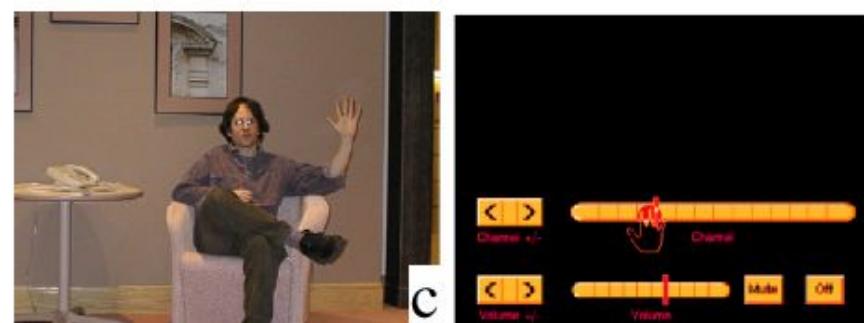
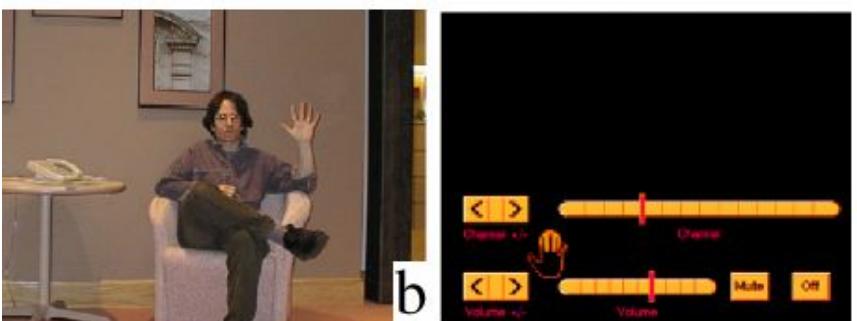


Figure from “Computer Vision for Interactive Computer Graphics,” W.Freeman et al, IEEE Computer Graphics and Applications, 1998 copyright 1998, IEEE

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

Convolution vs. (Cross) Correlation

- A **convolution** is an integral that expresses the amount of overlap of one function as it is shifted over another function.
 - convolution is a filtering operation
- **Correlation** compares the *similarity of two sets of data*. Correlation computes a measure of similarity of two input signals as they are shifted by one another. The correlation result reaches a maximum at the time when the two signals match best .
 - correlation is a measure of relatedness of two signals

Two operations are very similar, why convolution?

Convolution is correlation with pre-flipping of h.
Seems more cumbersome, why bother?

Two operations are very similar, why convolution?

Convolution is correlation with pre-flipping of h .
Seems more cumbersome, why bother?

- Correlation (unlike conv) is:
 - NOT commutative
(eg. $[1 \ 2 \ 3] \ ** \ [0 \ 1 \ 0] \neq [0 \ 1 \ 0] \ ** \ [1 \ 2 \ 3]$)
 - NOT associative: $a** (b**c) \neq (a**b)**c$ (in general)

Therefore, convolution is much more convenient for mathematical construction + close relation to Fourier transform.

What we have learned today?

- Image sampling and quantization
- Image histograms
- Images as functions
- Linear systems (filters)
- Convolution and correlation

Optional reading

Sections 3.1 and 3.2 of Richard Szeliski's excellent book (<https://szeliski.org/Book/>).