

# Lecture:

# Visual Bag of Words

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# What we will learn today

- Visual bag of words (BoW)
- Spatial Pyramid Matching
- Naive Bayes

# What we will learn today

- Visual bag of words (BoW)
- Spatial Pyramid Matching
- Naïve Bayes



# Bag of Words Models

Adapted from slides by Rob Fergus and  
Svetlana Lazebnik

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Object




# Bag of 'words'



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Origin 1: Texture Recognition







Example textures (from Wikipedia)



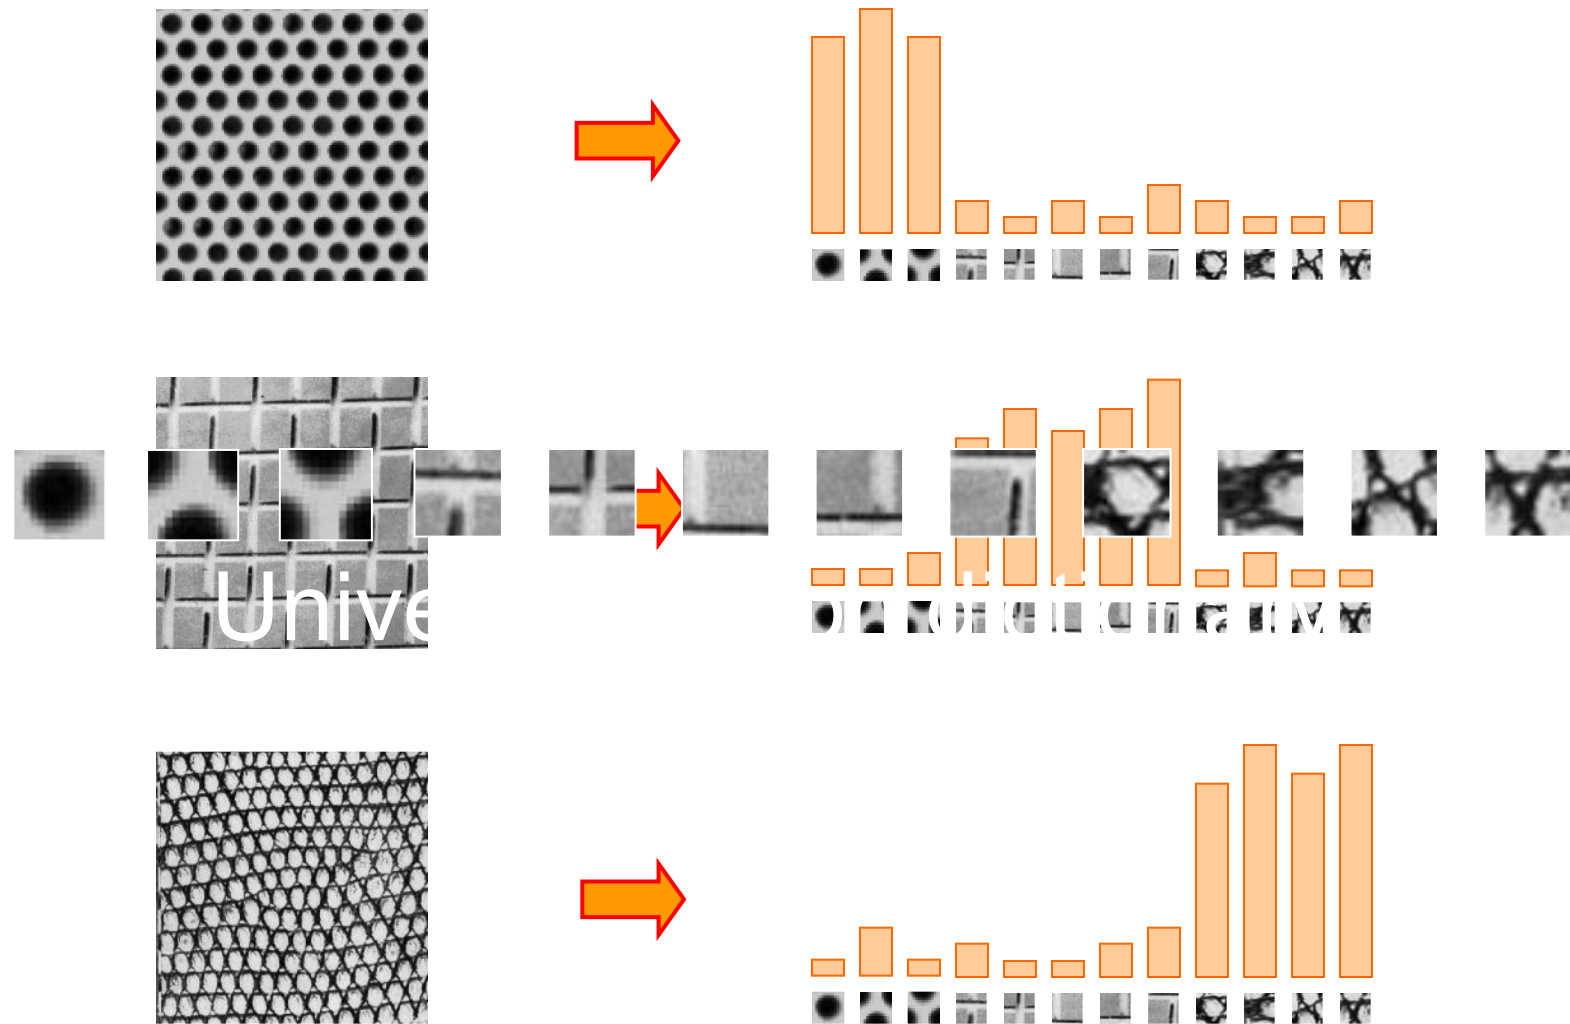
# Origin 1: Texture recognition

- Texture is characterized by the repetition of basic elements or *textons*



Julesz, 1981; Cula & Dana, 2001; Leung & Malik 2001; Mori, Belongie & Malik, 2001; Schmid 2001; Varma & Zisserman, 2002, 2003; Lazebnik, Schmid & Ponce, 2003

# Origin 1: Texture recognition



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna



# Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)

# Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)



US Presidential Speeches Tag Cloud  
<http://chir.ag/phernalia/preztags/>

# Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)



US Presidential Speeches Tag Cloud  
<http://chir.ag/phernalia/preztags/>

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)




US Presidential Speeches Tag Cloud

Adapted from slides by Juan Carlos Niechajew <http://chir.ag/phernalia/preztags/>



# Bags of features for object recognition




face, flowers, building

- Works pretty well for image-level classification and for recognizing object *instances*

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Bags of features for object recognition



class	bag of features	bag of features	Parts-and-shape model
	Zhang et al. (2005)	Willamowski et al. (2004)	Fergus et al. (2003)
airplanes	<b>98.8</b>	97.1	90.2
cars (rear)	98.3	<b>98.6</b>	90.3
cars (side)	<b>95.0</b>	87.3	88.5
faces	<b>100</b>	99.3	96.4
motorbikes	<b>98.5</b>	98.0	92.5
spotted cats	<b>97.0</b>	—	90.0

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna


# Bag of features

- First, take a bunch of images, extract features, and build up a “dictionary” or “visual vocabulary” – a list of common features
- Given a new image, extract features and build a histogram – for each feature, find the closest visual word in the dictionary




# Bag of features: outline

## 1. Extract features



# Bag of features: outline

1. Extract features
2. Learn “visual vocabulary”




# Bag of features: outline

1. Extract features
2. Learn “visual vocabulary”
3. Quantize features using visual vocabulary

# Bag of features: outline


1. Extract features
2. Learn “visual vocabulary”
3. Quantize features using visual vocabulary
4. Represent images by frequencies of “visual words”



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# 1. Feature extraction


- Regular grid
  - Vogel & Schiele, 2003
  - Fei-Fei & Perona, 2005



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# 1. Feature extraction

- Regular grid
  - Vogel & Schiele, 2003
  - Fei-Fei & Perona, 2005
- Interest point detector
  - Csurka et al. 2004
  - Fei-Fei & Perona, 2005
  - Sivic et al. 2005




Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# 1. Feature extraction

- Regular grid
  - Vogel & Schiele, 2003
  - Fei-Fei & Perona, 2005
- Interest point detector
  - Csurka et al. 2004
  - Fei-Fei & Perona, 2005
  - Sivic et al. 2005
- Other methods
  - Random sampling (Vidal-Naquet & Ullman, 2002)
  - Segmentation-based patches (Barnard et al. 2003)






## 2. Learning the visual vocabulary



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna



## 2. Learning the visual vocabulary




Clustering

Adapted from slides by Juan Carlos Nieves, and Ranjay Krishna

## 2. Learning the visual vocabulary



Visual vocabulary



Clustering

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# K-means clustering recap

- Want to minimize sum of squared Euclidean distances between points  $x_i$  and their nearest cluster centers  $m_k$


$$D(X, M) = \sum_{\text{cluster } k} \sum_{\substack{\text{point } i \text{ in} \\ \text{cluster } k}} (x_i - m_k)^2$$

- Algorithm:
- Randomly initialize K cluster centers
- Iterate until convergence:
  - Assign each data point to the nearest center
  - Recompute each cluster center as the mean of all points assigned to it

# From clustering to vector quantization



- Clustering is a common method for learning a visual vocabulary or codebook
  - Unsupervised learning process
  - Each cluster center produced by k-means becomes a codevector
  - Codebook can be learned on separate training set
  - Provided the training set is sufficiently representative, the codebook will be “universal”
- The codebook is used for quantizing features
  - A *vector quantizer* takes a feature vector and maps it to the index of the nearest codevector in a codebook
  - Codebook = visual vocabulary
  - Codevector = visual word

# Example visual vocabulary



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Image patch examples of visual words




Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna




# Visual vocabularies: Issues

- How to choose vocabulary size?
  - Too small: visual words not representative of all patches
  - Too large: quantization artifacts, overfitting
- Computational efficiency
  - Vocabulary trees (Nister & Stewenius, 2006)




### 3. Image representation



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Image classification

- Given the bag-of-features representations of images from different classes, how do we learn a model for distinguishing them?



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Uses of BoW representation

- Treat as feature vector for standard classifier
  - e.g k-nearest neighbors, support vector machine
- Cluster BoW vectors over image collection
  - Discover visual themes

# Large-scale image matching

- Bag-of-words models have been useful in matching an image to a large database of object *instances*



11,400 images of game covers  
(Caltech games dataset)



how do I find this image in the database?

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna



# Large-scale image search



## Build the database:

- Extract features from the database images
- Learn a vocabulary using k-means (typical k: 100,000)
- Compute *weights* for each word
- Create an inverted file mapping words  $\square$  images

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Weighting the words

- Just as with text, some visual words are more discriminative than others

***the, and, or***      vs.      ***cow, AT&T, Cher***

- the bigger fraction of the documents a word appears in, the less useful it is for matching
  - e.g., a word that appears in *all* documents is not helping us



# TF-IDF weighting

- Instead of computing a regular histogram distance, we'll weight each word by its *inverse document frequency*
- inverse document frequency (IDF) of word  $j$  =

$$\log \frac{\text{number of documents}}{\text{number of documents in which } j \text{ appears}}$$

# TF-IDF weighting

- To compute the value of bin  $j$  in image  $l$ :

*term frequency of  $j$  in  $l$*  **X** *inverse document frequency of  $j$*

# Inverted file

- Each image has ~1,000 features
- We have ~100,000 visual words
  - Each histogram is extremely sparse (mostly zeros)
- Inverted file
  - mapping from words to documents

```
"a": {2}
"banana": {2}
"is": {0, 1, 2}
"it": {0, 1, 2}
"what": {0, 1}
```

# Inverted file

- Can quickly use the inverted file to compute similarity between a new image and all the images in the database
  - Only consider database images whose bins overlap the query image

# Large-scale image search

query image



top 6 results



- Cons:
  - performance degrades as the database grows

Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# Large-scale image search

- Pros:
  - Works well for CD covers, movie posters
  - Real-time performance possible



real-time retrieval from a database of 40,000 CD covers

Adapted from slides by Nister & Stewenius, **Scalable Recognition with a Vocabulary Tree**

# Example bag-of-words matches



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna




# Example bag-of-words matches



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna

# What about spatial info?



Adapted from slides by Juan Carlos Niebles, and Ranjay Krishna